

AGENTICRAG-R1: Agentic Reinforcement Learning with Stack Memory for Multi-Step Reasoning, Retrieval and Memorizing

Xinke Jiang^{1,2,3,*}, Yue Fang^{1,2,3,*}, Zhibang Yang^{1,2,3,*}, Jiaran Gao^{1,2,3,*}, Zhixin Zhang^{1,2,3,*}, Tao Feng¹,
Rihong Qiu^{1,2,3}, Wentao Zhang¹, Hongxin Ding^{1,2}, Ruizhe Zhang^{1,2,3}, Yongxin Xu^{1,3},
Yuheng Huang⁴, Xu Chu^{2,3,5,†}, Junfeng Zhao^{2,3,†}, Yasha Wang^{1,6,†}

¹National Engineering Research Center of Software Engineering, Peking University, Beijing, China

²School of Computer Science, Peking University, Beijing, China

³Key Laboratory of High Confidence Software Technologies, Ministry of Education, Beijing, China

⁴GRG Banking Equipment Co., Ltd., Guangzhou, China

⁵Center on Frontiers of Computing Studies, Peking University, Beijing, China

⁶Peking University Information Technology Institute (Tianjin Binhai), Tianjin, China

*Equal contribution †Corresponding authors

{xinkejiang, yangzb}@stu.pku.edu.cn; {chu_xu, zhaojf, wangyasha}@pku.edu.cn

Abstract

Retrieval-Augmented Generation (RAG) improves the factuality of large language models (LLMs), yet existing RAG systems often struggle with complex, multi-step reasoning that requires *adaptive retrieval and continuous revision of intermediate contexts*. Recent reinforcement learning (RL)-based agentic RAG methods partially alleviate this issue, but typically rely on *coarse-grained action spaces* and *trajectory-level rewards*, resulting in weak reward assignment and a bias toward short-horizon, stereotyped reasoning template. To address, we propose AGENTICRAG-R1, a RL framework that deeply integrates reasoning, retrieval, and memory via a memory stack and fine-grained action space, supported by hierarchical action-aware rewards and an information-aware trajectory rejection strategy to enable effective long-horizon learning. Experiments across a diverse set of multi-hop, open-domain, and agentic reasoning benchmarks, spanning multiple backbone model sizes, demonstrate that AGENTICRAG-R1 consistently outperforms strong baselines. Moreover, AGENTICRAG-R1 learns more robust, interpretable, and memory-aware reasoning behaviors, highlighting the effect of fine-grained action modeling and information-aware optimization for long-horizon reasoning. Our code is anonymous available at <https://github.com/jiangxinke/Harness-RL/tree/AgenticRAG-R1-Whitebox>.

1 Introduction

Large Language Models (LLMs) have achieved impressive performance across a broad range of natural language processing tasks (Kaplan et al., 2020; Yang et al., 2024b; Guo et al., 2025; Vu




Figure 1: Conventional Agentic vs AGENTICRAG-R1 for multi-step reasoning and retrieval. Conventional methods adopt coarse-grained, append-only memory and shallow retrieval, leading to error accumulation and short-horizon. AGENTICRAG-R1 employs fine-grained, reversible memory actions and information-aware rollout selection, for robust long-horizon reasoning.

et al., 2024; Chang et al., 2024). However, despite their strengths, LLMs still remain prone to factual errors (Ji et al., 2023a; Cao et al., 2020; Ji et al., 2023b; Jiang et al., 2024b), outdated knowledge (He et al., 2022) and limited domain-specific expertise (Kandpal et al., 2023), making **Retrieval-Augmented Generation (RAG)** (Edge et al., 2024; Asai et al., 2024; Yang et al., 2024c) a promising solution by enabling access to external knowledge for improved accuracy and adaptability.

In contrast, many real-world decision-making tasks are inherently complicated (Gao et al., 2025), requiring models to decompose problems into multiple interdependent steps rather than relying on one-off reasoning or single-step retrieval (Shao et al., 2023; Press et al., 2022). Such settings

demand not only access to external knowledge, but also the ability to dynamically decide *when* and *what* to retrieve as the reasoning process unfolds (Lee et al., 2024; Jeong et al., 2024). Crucially, retrieved information could be continuously integrated into the model’s evolving internal state, influencing subsequent reasoning and decision making. This tight coupling between reasoning and retrieval makes their *deep integration* a critical yet underexplored challenge.

Early attempts to integrate reasoning and retrieval primarily relied on manually crafted prompting strategies (Press et al., 2022; Shao et al., 2023; Trivedi et al., 2023; Jiang et al., 2024a; Li et al., 2025), and were later extended through Supervised Fine-Tuning (SFT) to better coordinate their interaction (Wang et al., 2025b). While effective in specific settings, these approaches depend heavily on heuristic design and often overfit to **fixed reasoning patterns**, without **intrinsically** improving the model’s ability to jointly reason and retrieve (Chu et al., 2025a). More recently, Reinforcement Learning (RL) has emerged as a promising paradigm for enhancing complex reasoning, as demonstrated by OpenAI O1 (OpenAI, 2025) and DeepSeek-R1 (Guo et al., 2025). Motivated by this success, several studies have begun to apply RL to retrieval-augmented settings, training models to invoke external retrieval tools during reasoning (Jin et al., 2025; Dong et al., 2025a,b; Zhou et al., 2025; Liu et al., 2026). As illustrated in Figure 1, most existing RL-based RAG rely on **coarse-grained action modeling**, treating retrieval as a single or shallow action, which limits fine-grained control over complex tasks and biases policy learning toward **short-horizon**, easily **reproducible reasoning templates** than long-horizon retrieval and reasoning, posing two challenges:

❶ *Challenge 1. How to expose fine-grained, memory-aware action-making and reward allocation in multi-step reasoning and retrieval.* Deep integration of reasoning and retrieval requires LLMs not only to acquire **external information**, but also to continuously **manage and revise intermediate content memory** under noisy and uncertain conditions. In multi-step reasoning, both external retrieval and internal inference inevitably introduce noise: retrieved documents may be irrelevant or misleading (Yao et al., 2022b; Shinn et al., 2023; Liu et al., 2023), while intermediate reasoning steps may contain logical errors or overconfident assumptions (Jiang et al., 2024a; Xiong et al.,

2024; Hao et al., 2023). If such noisy intermediate conclusions are written into context memory without the ability to be inspected, revised, or discarded, they tend to propagate and amplify over steps, effectively corrupting the agent’s internal decision state (Ning et al., 2025; Hoscilowicz et al., 2024; Lai et al., 2024). However, most existing RL-based agentic methods adopt a highly **coarse-grained** action formulation, collapsing the entire reasoning and retrieval processes into a single latent and memory-appending <think> and <search> action template (Dong et al., 2025a,b; Zhou et al., 2025), which are then stored as irreversible content. This limitation is further exacerbated by reward design, where outcome rewards are evaluated only on final correctness, overlooking the contributions of intermediate actions throughout the trajectory (Shao et al., 2024; Jin et al., 2025). Without fine-grained, memory-aware action and reward signals, RL suffers from weak reward assignment and is biased toward **shallow, template-like strategies**.

❷ *Challenge 2. How to expose reinforcement learning to sufficiently diverse and high-information reasoning and retrieval rollouts.* Even with fine-grained action spaces and memory-aware reward mechanisms, RL may still fail to learn robust long-horizon reasoning if the training distribution is dominated by **low information trajectories**. Most multi-hop and retrieval-augmented tasks are simple—requiring only a **single retrieval step or shallow reasoning**—so these short-horizon trajectories often yield higher immediate rewards than more complex, long-horizon ones (Gao et al., 2025; Zhou et al., 2025). As a result, RL tends to overfit to **low-information**, trivially rewarding trajectories, rarely exercising actions such as revising, summarizing, or planning, and under-optimizing the long-horizon combinations of cognitive actions (Wang et al., 2025c).

To address the above challenges, we propose AGENTICRAG-R1, a reinforcement learning framework that achieves deep integration of reasoning, retrieval and memorizing through fine-grained action modeling and action-aligned reward mechanisms. ❶ For *CI*, we design a structured multi-action space together with a stack-based memory which represents the agent’s internal state across multi-step reasoning. The memory stack allows intermediate actions to be selectively pushed, revised, or popped, enabling the agent to inspect and correct noisy information accumulated during reasoning and retrieval. Specifically, our

framework supports distinct cognitive actions such as <think>, <search>, and <plan>, as well as memory-oriented actions such as <summary> and <backtrack>, which operate over the memory stack to refine or retract intermediate actions. This action design is further complemented by an action-aware reward attribution mechanism that assigns process rewards to each actions, providing fine-grained reward supervision. ② For *C2*, we introduce a dynamic training trajectories rejection strategy that rejects low-information trajectories and prioritizes high-reward rollouts while up-weighting inputs whose rollouts exhibit meaningful dispersion. Inspired by rejection sampling (Liu et al.; Xiong et al., 2025), we define a **cross-batch, information-awareness metric (Upper Confidence Bound-like)** (Russo and Roy, 2014; Russo and Van Roy, 2013) that evaluates each trajectory based on both its *rollout reward* and the *sample variance*. This approach ensures trajectories eliciting diverse reasoning behaviors and requiring multi-step, long-horizon decision-making receive higher priority during training. In summary, our main contributions are as follows:

- We first formalize the agentic reinforcement learning process within a search-and-memory environment, and introduce a fine-grained, memory-aware action space together with action-aware process-level rewards. This design enables explicit supervision of heterogeneous cognitive actions and effectively mitigates error accumulation in the agent’s memory (*C1*).
- We introduce a dynamic, information-aware trajectories rejection strategy that filters low-information trajectories and prioritizes challenging, long-horizon samples with high behavioral variance, ensuring RL effectively explores and optimizes complex, multi-step reasoning strategies (*C2*).
- We evaluate AGENTICRAG-R1 on multiple multi-hop and RAG benchmarks, demonstrating improved reasoning performance, generalization, and applicability to complex downstream tasks, i.e., report generation.

2 AGENTICRAG-R1

In this section, we present AGENTICRAG-R1, as illustrated in Figure 2, which consists of two key components: ① action-centric training paradigm that supports multi-turn, multi-action rollouts with reasoning, retrieval, and memory operations, and

incorporates hierarchical, action-aware reward modeling for fine-grained, token-level reward assignment; ② Information-Gain-Aware dynamic rollout selection strategy, improves long-horizon exploration and data efficiency during training. More details of Notation Table, Full Algorithm, Prompt, Proofs, are in Appendix C, B, G, K.

2.1 Action Modeling with Hierarchical Rewards

Memory-aware Multi-Action Modeling. To overcome *C1*, we introduce a structured *multi-action mechanism* integrated with the last-in-first-out (LIFO) memory stack that supports push and pop operations. Each stack element corresponds to a discrete and semantically meaningful reasoning unit, serving as the basic execution primitive of the agent. The memory stack enables step-wise **action decomposition** of reasoning trajectories, supports dynamic rollback through stack manipulation, and maintains a reversible and interpretable record of the agent’s internal decision process:

- **<Plan>** (push): Generate and store a sub-goal, task decomposition, or high-level strategy for long-horizon reasoning.
- **<Think>** (push): Record internal reasoning, such as logical deduction or decision-making, to support subsequent inference.
- **<Search>** (push): Formulate a query with structured parameters and invoke the retrieval module. This action explicitly determines what and how external knowledge is accessed by external tools.
- **<Backtrack>** (pop, then push): Remove one or more top elements upon detecting inconsistencies or errors to restore a prior valid reasoning state for alternative exploration.
- **<Summary>** (pop, then push): Compress recent memory entries into a concise representation and push the summary back onto the stack, for context management and information preservation.
- **<Conclusion>** (push): Indicate that the current state represents a candidate final answer, concluding the reasoning sequence.

Hierarchical Reward Modeling To better model interaction outcomes arising from multi-turn reasoning and multi-action execution—particularly those involving **external retrieval engines and the memory stack**—we introduce a *hierarchical reward modeling* framework. Specifically, we decompose the overall reward signal into two parts:




Figure 2: Overall Framework of AGENTICRAG-R1.

a *outcome reward* and a *process reward*. The outcome reward reflects the quality of the final outcome and is uniformly distributed across valid generation tokens, while the process reward provides fine-grained supervision and is assigned to tokens for corresponding actions.

Formally, for a rollout indexed by k , let $y^{(k)} = (x_1^{(k)}, \dots, x_{|y^{(k)}|}^{(k)})$ denote the generated token sequence of the k -th trajectory. We define per-token reward as:

$$r_i^{(k)} = (1 - \mathbb{I}_{\text{mask},i}^{(k)}) r_{\text{out}}^{(k)} + \mathbb{I}_{\text{act},i}^{(k)} r_{\text{proc},i}^{(k)}, \quad i \in [1, |y^{(k)}|], \quad (1)$$

where $\mathbb{I}_{\text{mask},i}^{(k)}$ indicates whether token $x_i^{(k)}$ is masked from policy optimization (e.g., retrieved observations) to avoid those false gradient updates, and $\mathbb{I}_{\text{act},i}^{(k)}$ indicates whether the token corresponds to this action invocation. This hierarchical reward formulation allows outcome-level supervision to be shared across the entire reasoning trajectory, while process-level rewards are selectively applied to action tokens, providing fine-grained control over decisions:

- **Outcome Reward** $r_{\text{out}}^{(k)}$: evaluates the quality of the final answer produced by the trajectory. We extract the predicted answer from `<conclusion>` $a_{\text{conclusion}}^{(k)}$ from $y^{(k)}$ and compare it with the ground-truth a_{gold} via Exact Match (EM) for discrete answers and F1 score for graded answers:

$$r_{\text{out}}^{(k)} = \text{Acc}(a_{\text{conclusion}}^{(k)}, a_{\text{gold}}) \in [0, 1]. \quad (2)$$

- **Process Rewards** $r_{\text{proc},i}^{(k)}$: Beyond the final outcome, we introduce action-conditioned process

rewards to supervise intermediate decisions during reasoning. These rewards are applied only at action tokens and encourage syntactically valid actions, effective external retrieval, and robust memory manipulation. Specifically, the process reward is defined as:

$$r_{\text{proc},i}^{(k)} = r_{\text{fmt},i}^{(k)} + r_{\text{rag},i}^{(k)} + r_{\text{mem},i}^{(k)}, \quad (3)$$

which contains the three types:

- **① Action Format Reward** ($r_{\text{fmt},i}^{(k)}$), enforces syntactic correctness at action level. Let $\mathbb{I}_{\text{syntax}}(a_i)$ be an indicator that equals 1 if action a_i is properly formatted with valid `<ACTION>...</ACTION>` tags, and 0 otherwise. $r_{\text{fmt},i}^{(k)}$ is defined as

$$r_{\text{fmt},i}^{(k)} = \alpha_{\text{type}}(a_i), \quad \text{if } \mathbb{I}_{\text{syntax}}(a_i) = 1, \quad (4)$$

where $\alpha_{\text{type}}(a_i) = 1.0$ for CONCLUSION actions and 0.5 for other valid action types, otherwise 0.

- **② Search-based Reward** ($r_{\text{rag},i}^{(k)}$). To encourage effective external retrieval, we assign a reward based on semantic relevance of retrieved content. Upon a `<Search>` action, a large-parameterized *general reward model* (GRM) (Wang et al., 2025a) evaluates the alignment between the issued query, the retrieved evidence, and the current task context, producing a relevance score:

$$r_{\text{rag},i}^{(k)} = \text{GRM}(q_i^s, o_i, q) \in [0, 1], \quad (5)$$

where q_i^s denotes the search query, o_i the retrieved observation, and q the original user input.

- **Memory-aware Reward** ($r_{\text{mem},i}^{(k)}$) supervises *useful* memory operations rather than their mere occurrence. For memory-related actions (Backtrack, Summary), we also employ the GRM to assess whether the operation is *reasonable* and *beneficial* given the current reasoning context and memory state. Specifically, we define:

$$r_{\text{mem},i}^{(k)} = \text{GRM}(a_i, M_i, q) \in [0, 1], \quad (6)$$

where M_i denotes the current memory stack. We use separate fixed prompts for retrieval and memory evaluation, as detailed in Appendix G.

Reward Budget and Proportional Normalization. To prevent reward hacking via excessive action invocation, we impose a per-trajectory budget on process rewards. Let $r_{\text{proc}}^{(k)} = \sum_i r_{\text{proc},i}^{(k)}$ denote the accumulated process reward. Process reward is normalized as:

$$\tilde{r}_{\text{proc},i}^{(k)} = \begin{cases} r_{\text{proc},i}^{(k)}, & R_{\text{proc}}^{(k)} \leq B_{\text{proc}}, \\ \frac{B_{\text{proc}}}{R_{\text{proc}}^{(k)}} r_{\text{proc},i}^{(k)}, & R_{\text{proc}}^{(k)} > B_{\text{proc}}, \end{cases} \quad (7)$$

where B_{proc} is a fixed reward budget. This ensures that when the total process reward exceeds the budget, all process rewards are proportionally scaled to share the same budget, thereby discouraging excessive or redundant action invocation.

2.2 Information-Aware Trajectories Rejection

Despite fine-grained, action-level reward, **C2** still remains: how to distinguish *informative* trajectories from noisy or redundant ones when multiple stochastic rollouts are generated per input. In practice, rollouts sampled from the same prompt can exhibit substantially different reasoning paths and reward outcomes, while rollouts across different inputs vary widely in their overall learning utility. Naively treating all rollouts equally during optimization can dilute gradient signals and bias training toward **short-sighted, low-effort, or stereotyped templated behaviors**. To address, we propose an *Information-Aware Trajectories Rejection* mechanism that explicitly accounts for both *input-level uncertainty* and *rollout-level quality*. Unlike prior approaches that perform selection independently per input, our method operates *globally* over a mini-batch of rollouts, enabling competition across both inputs and trajectories.

Batch-level Rollout Ensemble. At each training step, we sample a mini-batch of inputs $\mathcal{B} = \{q^{(1)}, \dots, q^{(B)}\}$. For each input $q^{(b)}$, the policy π_θ generates K stochastic rollouts $y^{(b)} = \{y^{(b,1)}, \dots, y^{(b,K)}\}$, where $y^{(b,i)} \sim \pi_\theta(\cdot \mid q^{(b)})$. All rollouts in the mini-batch form a unified rollout set $\mathcal{Y}_\mathcal{B} = \bigcup_{b=1}^B y^{(b)}$ with cardinality $|\mathcal{Y}_\mathcal{B}| = B \times K$. Each rollout $y^{(b,i)}$ is assigned with a scalar reward $r^{(b,i)}$. For each input $q^{(b)}$, we compute the empirical reward mean and variance across its rollouts to estimate expected performance:

$$\mu^{(b)} = \frac{1}{K} \sum_{i=1}^K r^{(b,i)}, \quad (\sigma^{(b)})^2 = \frac{1}{K} \sum_{i=1}^K \left(r^{(b,i)} - \mu^{(b)} \right)^2,$$

Here, the variance $(\sigma^{(b)})^2$ provides a lightweight measure of *input-level informativeness*, reflecting the variability of learning signals across reasoning and retrieval trajectories for a given input.

Information-Aware Rejection. Given the batch-level rollout set $\mathcal{Y}_\mathcal{B} = \{y^{(b,i)}\}_{b=1, i=1}^{B,K}$, we assign each rollout an *information-aware score* that jointly accounts for *input-level diversity* and *rollout-level quality*. Concretely, for i -th rollout of input $x^{(b)}$, we define it as Upper Confidence Bound (UCB) following (Russo and Roy, 2014; Russo and Van Roy, 2013):

$$\tilde{r}^{(b,i)} = (\sigma^{(b)})^2 \cdot r^{(b,i)} + \mu^{(b)}, \quad (8)$$

where $r^{(b,i)}$ denotes the aggregated scalar reward of the rollout, and $\mu^{(b)}$ and $(\sigma^{(b)})^2$ are the mean and variance of rollout rewards under the same input $x^{(b)}$. This scoring function captures two complementary signals. The variance term $(\sigma^{(b)})^2$ reflects how much the learning signal varies across different reasoning and retrieval trajectories induced by the same input, and thus emphasizes inputs that admit diverse, non-trivial solution paths. The rollout-level reward $r^{(b,i)}$ distinguishes higher-quality trajectories within the same input context. By multiplicatively coupling these terms, the score selectively amplifies **high-reward rollouts when meaningful diversity exists**, while suppressing uninformative or degenerate trajectories. Additive mean term $\mu^{(b)}$ serves as a stabilizing baseline to prevent overly rejection.

Rollout rejection is then performed *globally* over the batch-level rollout set. Specifically, we rank all rollouts $y^{(b,i)} \in \mathcal{Y}_\mathcal{B}$ according to $\tilde{r}^{(b,i)}$ and retain only the top fraction $p \in (0, 1]$. Rollouts not

selected are discarded and do not participate in policy optimization. By performing rejection over the joint $batch \times rollout$ space, this mechanism prioritizes trajectories that are both locally high-quality and globally informative, encouraging exploration of diverse reasoning strategies while improving long-horizon reasoning.

3 Experiments

We conduct extensive experiments to evaluate AGENTICRAG-R1 and address the following key research questions:

- **RQ1:** Does AGENTICRAG-R1 consistently outperform strong baselines on in-domain and out-of-domain QA?
- **RQ2:** How do memory-aware actions and trajectory rejection jointly affect overall model performance?
- **RQ3:** What is the impact of retrieval- and memory-based rewards on the quality of learning and reasoning?
- **RQ4:** How sensitive is AGENTICRAG-R1 to the chosen rejection threshold in trajectory filtering?
- **RQ5:** How does AGENTICRAG-R1 effectively handle long-horizon reasoning with larger step budgets?

3.1 Experimental Setup

❶ **Training Data.** We train our models and baselines on a curated multi-hop question answering dataset constructed from the training splits of HotpotQA (Yang et al., 2018) and 2WikiMultiHopQA (Ho et al., 2020). To focus on genuinely non-trivial reasoning scenarios, we filter out instances that require no external retrieval or can be solved with only a single, trivial retrieval step. Detailed dataset construction procedures and statistics are provided in Appendix D. ❷ **Evaluation Benchmarks.** We evaluate our method on ten benchmarks spanning three settings: *multi-hop QA* (HotpotQA, 2Wiki, MusiQue (Trivedi et al., 2022), Bamboogle (Press et al., 2022)), *open-domain QA* (Natural Questions (Kwiatkowski et al., 2019), TriviaQA (Joshi et al., 2017)), and *agentic benchmarks* (FRAMES (Krishna et al., 2025)). Additional benchmark details are reported in Appendix D. ❸ **Backbone Models.** We adopt Qwen2.5 (Yang et al., 2024a) as the backbone language model across all experiments, with parameter scales of 1.5B, 3B, and 7B. ❹ **RAG Tools.** We employ a Wikipedia-based search tool built

on a dump dated November 1, 2023. ❺ **Baselines.** We compare our approach against a broad range of baselines spanning *No-RAG*, *Naive RAG*, *Agentic RAG*, and *RL-based Agentic* paradigms. These include Base, CoT (Wei et al., 2022a), FS-RAG (Trivedi et al., 2023), FL-RAG (Khandelwal et al., 2020), ReAct (Yao et al., 2022b), IR-CoT (Trivedi et al., 2023), TC-RAG (Jiang et al., 2024a), ReSearch (Chen et al., 2025), Search-R1 (Jin et al., 2025), AEPO (Dong et al., 2025a), ARPO (Dong et al., 2025b), and Mem1 (Zhou et al., 2025). Baseline descriptions are provided in Appendix E. ❻ **Evaluation Metrics.** We report the **F1 score** (Yacoub and Axman, 2020) as the primary evaluation metric for all question-answering benchmarks (Appendix F). ❼ **Additional Experiments.** More experiments are listed in Appendix I.

3.2 Main Result Analysis

To address **RQ1**, we evaluate AGENTICRAG-R1 on two types of tasks using three backbones (Table 1) and analyze its training behaviors and learning dynamics (Figure 3). Based on these analyses, we highlight the following main observations:

❶ **Strong Performance Compared with Baselines.** Across in-domain and out-of-domain multi-hop QA benchmarks (Table 1), AGENTICRAG-R1 achieves the best or near-best performance across backbone model sizes. On the in-domain datasets, AGENTICRAG-R1 consistently surpasses the strongest baselines, with, for example, absolute gains of **+3.02%** on 2Wiki and **+6.76%** on HotpotQA using the 3B backbone. On out-of-domain benchmarks, AGENTICRAG-R1 exhibits strong generalization, e.g., achieving an absolute improvement of about **+7.2%** on TriviaQA with the 1.5B backbone and **+7.1%** on Bamboogle with the 7B backbone over the strongest Agentic RAG baselines. Notably, these gains hold across backbone scales, demonstrating the effectiveness and robustness of our reinforcement learning-based optimization in this setting.

❷ **Learning Dynamics of Cognitive Actions and Information-Aware Rejection.** We examine the evolution of cognitive actions during training and the associated learning dynamics, corroborating trends using reward trajectories from Figure 5. Using the Qwen2.5-3B model as a primary example, <think> actions steadily increase, reflecting a growing reliance on internal reasoning. Actions related to intermediate reasoning, namely <summary> and <backtrack>, fluctuate

Method		In-Domain		Out-of-Domain					
Paradigm	Approach	2Wiki	HotpotQA	Bamboogle	FRAMES	MusiQue	NQ	TriviaQA	Avg
Qwen2.5-1.5B									
No RAG	Base COT	21.95	20.20	5.87	9.51	8.26	11.31	32.81	15.70
		16.75	17.90	19.74	8.43	7.22	9.82	25.17	15.00
Naive RAG	FS-RAG FL-RAG	21.37	26.55	15.31	11.40	9.50	17.74	44.73	20.94
		26.08	28.14	14.84	12.58	10.30	21.94	48.59	23.21
Agentic RAG	ReACT	12.43	23.48	12.54	8.63	7.99	21.06	33.91	17.15
	IRCOT	21.79	26.68	17.49	9.32	8.73	22.93	44.19	21.59
	TCRAG	25.23	22.29	11.63	10.27	7.90	14.21	34.05	17.94
RL-based	ReSearch	25.55	30.45	22.16	9.21	9.54	29.07	44.65	24.38
	Search-R1	26.51	20.66	13.80	9.65	8.36	11.93	25.76	16.67
	AEPO	17.34	14.75	11.58	8.98	7.11	9.55	23.02	13.19
	ARPO	20.02	15.48	12.38	9.11	6.16	9.19	23.96	13.76
	Mem1	14.43	14.34	12.03	8.18	6.08	18.41	22.06	13.65
	Ours	29.53 (+3.02)	30.88 (+0.43)	20.00 (-2.16)	13.46 (+0.88)	16.48 (+6.18)	31.35 (+2.28)	51.83 (+3.24)	27.65 (+3.27)
Qwen2.5-3B									
No RAG	Base COT	23.98	24.08	9.45	8.01	9.70	14.27	38.84	18.33
		18.90	23.82	20.80	7.16	10.47	16.53	39.62	19.61
Naive RAG	FS-RAG FL-RAG	15.47	25.85	10.48	10.42	7.64	19.84	45.38	19.30
		16.80	26.78	11.05	9.19	7.29	21.93	48.50	20.22
Agentic RAG	ReACT	25.09	34.37	24.86	10.53	13.92	27.19	46.04	26.00
	IRCOT	15.89	24.50	25.27	6.79	12.43	27.86	49.19	23.13
	TCRAG	28.47	21.94	17.59	7.69	8.99	20.74	50.46	22.27
RL-based	ReSearch	27.23	33.96	15.09	10.00	9.47	34.61	53.93	26.33
	Search-R1	29.90	37.24	29.90	10.76	13.53	34.73	55.08	30.16
	AEPO	23.01	28.71	22.09	12.26	11.70	26.76	47.78	24.62
	ARPO	29.55	36.48	27.32	13.49	13.38	33.29	53.66	29.60
	Mem1	18.06	20.15	5.19	4.99	4.47	19.18	33.09	15.02
	Ours	32.92 (+3.02)	44.00 (+6.76)	31.48 (+1.58)	16.23 (+2.74)	16.48 (+2.56)	37.15 (+2.42)	56.62 (+1.54)	33.55 (+3.39)
Qwen2.5-72B									
No RAG	Base COT	25.41	26.63	17.86	12.52	12.15	19.72	49.08	23.34
		23.55	29.10	37.56	17.60	14.35	22.47	49.33	27.71
Naive RAG	FS-RAG FL-RAG	17.71	29.21	16.86	12.52	10.74	16.82	35.02	19.84
		19.78	34.42	24.10	12.10	12.46	19.72	42.66	23.61
Agentic RAG	ReACT	27.51	42.81	27.63	15.29	19.34	30.01	54.55	31.02
	IRCOT	36.45	26.29	21.90	6.78	8.39	19.63	49.43	24.12
	TCRAG	29.70	40.83	25.13	16.46	17.56	29.01	54.78	30.50
RL-based	ReSearch	30.03	30.39	30.42	15.61	12.58	23.69	48.25	27.28
	Search-R1	35.03	38.89	42.04	18.01	19.08	29.59	55.91	26.94
	AEPO	19.88	13.85	13.24	7.24	5.85	9.93	17.53	12.50
	ARPO	30.71	25.20	32.94	12.18	12.71	17.80	40.16	24.53
	Mem1	25.29	29.98	36.50	14.15	14.13	26.38	51.04	28.21
	Ours	38.34 (+1.89)	45.15 (+2.34)	49.21 (+7.17)	19.44 (+1.43)	22.01 (+2.67)	23.60 (-6.41)	58.45 (+2.54)	36.60 (+5.58)

Table 1: Performance comparison (%) on multi-hop, retrieval-augmented benchmarks, including 2Wiki, HotpotQA, Bamboogle, FRAMES, MusiQue, NQ, and TriviaQA. **Bold** indicates the best, and underline indicates the second.

but overall increase; these behaviors help manage memory and remove noise from previous contexts. In contrast, <search> remains relatively stable with minor increases. <conclusion> actions rise consistently, reducing cases where the model fails to produce a final answer. For the smaller Qwen2.5-1.5B model, similar trends are observed for <think>, <summary>, <backtrack>, and <plan>. However, <search> frequency is lower, likely due to limited context capacity: performing excessive retrievals can disrupt reasoning, so the model reduces search while still successfully completing reasoning episodes, as reflected by increased <conclusion>.

Additional training statistics track overall model behavior throughout training (Figure 5). Over 0–300 training steps, rewards steadily increase while

action entropy shows a general downward trend, suggesting convergence toward more consistent policies. Response length and number of reasoning turns initially rise and then decline, indicating that the model discovers more efficient reasoning and retrieval strategies over time, ultimately producing answers faster and more accurately. Collectively, these analyses reveal how AGENTICRAG-R1 adapts its cognitive actions and information-aware rejection throughout training, balancing exploration and exploitation to optimize multi-hop reasoning.

3.3 Component Analysis

We analyze the effects of key model components, process-level rewards, and rejection thresholds on performance across both in-domain and out-of-domain tasks.

Paradigm	Method Approach	In-Domain		Out-of-Domain					
		2Wiki	HotpotQA	Bamboogle	FRAMES	MusiQue	NQ	TriviaQA	Avg
Full Model	AGENTICRAG-R1	32.92	44.00	31.48	16.23	16.48	37.15	56.62	33.55
Component Ablation	w/o MemAction	27.54	28.77	24.60	9.83	13.86	35.10	51.91	27.37
	w/o Reject	26.86	28.33	27.69	12.98	12.68	23.67	45.80	25.43
	w/o MemAct&Reject	24.56	28.61	25.49	11.61	12.54	22.07	44.83	24.24
Reward Ablation	w/o r_{rag}	23.70	34.71	19.68	7.70	11.45	29.07	52.39	25.53
	w/o r_{mem}	27.49	33.78	17.14	7.41	7.63	22.00	55.29	24.39
	w/o r_{rag} & r_{mem}	22.98	26.42	15.74	5.33	6.07	18.62	44.40	19.94

Table 2: Ablation analysis of component and reward designs in AGENTICRAG-R1 on Qwen2.5-3B.




Figure 3: Action-level frequency distributions during RL training across training steps.

❶ Model Component Ablation. We evaluate key components of AGENTICRAG-R1 in Table 2 for **RQ2**. Removing memory-related actions (*w/o MemAction*) substantially degrades performance, e.g., 2Wiki drops from 32.92% to 27.54% and HotpotQA from 44.00% to 28.77%. Removing information-aware rejection (*w/o Reject*) also reduces the average score from 33.55% to 25.43%, while removing both components further lowers it to 24.24%, confirming their complementary contributions.

❷ Process-level Reward Ablation. Table 2 also analyzes reward design for **RQ3**. Removing the RAG reward (*w/o r_{rag}*) hurts 2Wiki (32.92% to 23.70%) and HotpotQA (44.00% to 34.71%), while removing the memory reward (*w/o r_{mem}*) weakens out-of-domain performance, e.g., MusiQue drops from 16.48% to 7.63%. Removing both rewards yields the largest drop (Avg 19.94%), showing that both signals are necessary for effective reasoning and memory utilization.

❸ Sensitivity to Rejection Threshold. We analyze the impact of varying the rejection threshold on model performance and training behavior (Figure 4) to address **RQ4**. Accuracy, measured by F1-Score, initially decreases and reaches a minimum around a threshold of 0.2, then steadily improves as the threshold increases, peaking at 0.6




Figure 4: Sensitivity analysis of Rejection threshold on Qwen2.5-3B. The line plot shows F1-Score, and the bar chart depicts the number of rejected training samples.

and 0.8. Correspondingly, the number of rejected training samples decreases with higher thresholds, from over 20,000 at 0.2 to the fewest at 0.8. This trend indicates that moderate to high thresholds effectively balance answer generation and rejection.

❹ Long-Horizon Experiments. To address **RQ5**, we progressively increase the maximum number of allowed steps from 10 to 15, 20, 25, and 30, thereby evaluating its capability to perform long-horizon reasoning effectively.

As shown in Table 3, our model exhibits clear and consistent gains as the maximum step budget increases, demonstrating strong scalability in long-horizon settings. This trend holds across all evaluated datasets, indicating that the model can ef-

Method	Max Steps				
	10	15	20	25	30
<i>2Wiki</i>					
ReAct	18.23	22.45	25.17	27.89	29.03
Search-R1	24.12	38.02	37.51	33.08	44.27
TC-RAG	22.78	37.35	36.92	32.41	43.10
AGENTICRAG-R1	23.91	31.81	32.53	38.61	47.45
<i>Bamboogle</i>					
ReAct	15.67	18.92	20.34	21.76	22.19
Search-R1	37.05	34.11	37.02	28.03	28.19
TC-RAG	36.42	33.34	36.36	27.27	27.23
AGENTICRAG-R1	10.03	31.81	14.29	21.43	35.47
<i>TriviaQA</i>					
ReAct	28.76	30.12	31.05	31.42	31.68
Search-R1	34.21	36.12	33.57	33.09	32.28
TC-RAG	33.57	35.48	32.91	32.44	31.62
AGENTICRAG-R1	37.53	32.53	39.45	43.50	55.28

Table 3: Long-horizon performance under different maximum step budgets across multiple datasets.

Variant	2Wiki	HotpotQA	Avg.
Full AGENTICRAG-R1	32.92	44.00	33.55
w/o <Plan>	31.00	41.80	31.60
w/o <Summary>	29.90	40.10	30.50
w/o <Backtrack>	28.80	38.60	29.70
w/o both memory actions	27.54	28.77	27.37

Table 4: Leave-one-action-out ablation of AGENTICRAG-R1 on Qwen2.5-3B.

effectively leverage extended interaction horizons to perform deeper multi-step reasoning and retrieval. In contrast, TC-RAG shows unstable or even degraded performance when the horizon grows, suggesting limited ability to accumulate and utilize information over long trajectories. These results highlight that simply increasing the allowed steps is insufficient; rather, the proposed design is crucial for unlocking effective long-horizon reasoning capabilities. Furthermore, compared to ReAct and Search-R1, which either plateau early or fluctuate inconsistently across step budgets, AGENTICRAG-R1 exhibits monotonically improving trends on most benchmarks, further validating its effectiveness in long-horizon settings.

⑤ Individual Action Ablation. Table 4 analyzes the contribution of individual actions. Removing <Plan> decreases the average F1 score by 1.95%, while removing <Summary> and <Backtrack> causes larger drops of 3.05% and 3.85%, respectively. Removing both memory actions yields the largest degradation (6.18%), demonstrating that these actions are complementary for error correction and context compression.

⑥ Effect of SFT Initialization. We compare SFT only, cold-start RL, and RL with SFT

initialization on Qwen2.5-3B. Their average F1 scores across seven benchmarks are $22.40 \pm 0.50\%$, $33.55 \pm 0.42\%$, and $35.20 \pm 0.40\%$, respectively. SFT alone underperforms RL, while SFT initialization further improves RL performance. These results indicate that SFT is not required for our framework, but stronger initialization remains beneficial.

4 Related Works

Reinforcement Learning for LLMs Reinforcement learning has recently shown strong effectiveness in improving LLM reasoning and tool-use capabilities, as demonstrated by OpenAI-O1 (OpenAI, 2025) and DeepSeek-R1 (Guo et al., 2025). While policy gradient methods such as PPO (Schulman et al., 2017) are widely used, they incur high computational costs due to repeated on-policy updates. Alternative approaches, including DPO (Rafailov et al., 2023) and GRPO (Shao et al., 2024), improve training efficiency but face challenges such as off-policy bias or limited exploration. Recent advances target multi-step reasoning and tool interaction, proposing techniques to stabilize training and improve exploration (Dong et al., 2025a,b; Zhou et al., 2025). Nevertheless, existing RL methods remain limited by coarse actions, sparse rewards, and error accumulation in long-horizon reasoning.

5 Conclusion and Future Work

We presented AGENTICRAG-R1, a reinforcement learning framework for agentic retrieval-augmented reasoning that integrates fine-grained agentic actions, structured memory control, and action-aware rewards. Experiments across multiple reasoning benchmarks show that AGENTICRAG-R1 consistently outperforms strong RAG and agentic baselines while supporting more effective long-horizon reasoning. Several challenges remain for future work. In particular, jointly optimizing high-level actions and their associated arguments, as well as resolving conflicts among different agentic actions, remains difficult and calls for more principled coordination and credit assignment. We also plan to evaluate AGENTICRAG-R1 in broader vertical domains and more open-ended scenarios, such as deep research and long-horizon tasks.

Acknowledgments

This work was supported by the National Natural Science Foundation of China (Grant No. 62506010).

Limitations

Despite the promising results, our framework does have some limitations that need to be addressed. First, the computational overhead introduced by the memory stack system and dynamic retrieval processes can be relatively high. While these components are essential for managing complex queries and minimizing error accumulation, they do contribute to increased computational costs. In large-scale or real-time applications, this could present challenges related to processing speed and efficiency. Although we have proposed solutions like speculative sampling to alleviate some of this overhead, further optimization is still needed to enhance performance and scalability in time-sensitive environments. Second, our training strategy relies on a rejection-based sampling mechanism that prioritizes more challenging examples to enhance reasoning ability. While effective, this approach requires careful tuning of several hyperparameters, as different configurations can lead to varying performance outcomes. This additional tuning effort may complicate practical deployment and affect performance when transferring the framework to new datasets or tasks. Developing more adaptive or self-calibrating sampling strategies to reduce reliance on manual hyperparameter selection is therefore an important direction for future work.

Ethical considerations

To evaluate the performance of the proposed method, we carried out experiments exclusively on publicly accessible benchmark datasets, such as GAIA, HotpotQA, MusiQue, and other widely used resources, in accordance with their respective licenses and usage policies. We further confirm that no personally identifiable information was used in this study, and that no human or animal subjects were involved in the research.

References

Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. 2024. Self-rag: Learning to retrieve, generate, and critique through self-reflection. In *ICLR*.

Meng Cao, Yue Dong, Jiapeng Wu, and Jackie Chi Kit Cheung. 2020. [Factual error correction for abstractive summarization models](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 6251–6258, Online. Association for Computational Linguistics.

Yupeng Chang, Xu Wang, Jindong Wang, Yuan Wu, Linyi Yang, Kaijie Zhu, Hao Chen, Xiaoyuan Yi, Cunxiang Wang, Yidong Wang, and 1 others. 2024. A survey on evaluation of large language models. *ACM transactions on intelligent systems and technology*, 15(3):1–45.

Mingyang Chen, Linzhuang Sun, Tianpeng Li, Haoze Sun, Yijie Zhou, Chenzheng Zhu, Haofen Wang, Jeff Z Pan, Wen Zhang, Huajun Chen, and 1 others. 2025. Learning to reason with search for llms via reinforcement learning. *arXiv preprint arXiv:2503.19470*.

Tianzhe Chu, Yuexiang Zhai, Jihan Yang, Shengbang Tong, Saining Xie, Dale Schuurmans, Quoc V Le, Sergey Levine, and Yi Ma. 2025a. Sft memorizes, rl generalizes: A comparative study of foundation model post-training. *arXiv preprint arXiv:2501.17161*.

Xu Chu, Xinke Jiang, Rihong Qiu, Jiaran Gao, and Junfeng Zhao. 2025b. Model shapley: Find your ideal parameter player via one gradient backpropagation. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*.

Guanting Dong, Licheng Bao, Zhongyuan Wang, Kangzhi Zhao, Xiaoxi Li, Jiajie Jin, Jinghan Yang, Hangyu Mao, Fuzheng Zhang, Kun Gai, and 1 others. 2025a. Agentic entropy-balanced policy optimization. *arXiv preprint arXiv:2510.14545*.

Guanting Dong, Hangyu Mao, Kai Ma, Licheng Bao, Yifei Chen, Zhongyuan Wang, Zhongxia Chen, Jiazhen Du, Huiyang Wang, Fuzheng Zhang, and 1 others. 2025b. Agentic reinforced policy optimization. *arXiv preprint arXiv:2507.19849*.

Mingxuan Du, Benfeng Xu, Chiwei Zhu, Xiaorui Wang, and Zhendong Mao. 2025. Deepresearch bench: A comprehensive benchmark for deep research agents. *arXiv preprint arXiv:2506.11763*.

Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, and Jonathan Larson. 2024. [From local to global: A graph rag approach to query-focused summarization](#). Preprint, arXiv:2404.16130.

Yuchen Fan, Kaiyan Zhang, Heng Zhou, Yuxin Zuo, Yanxu Chen, Yu Fu, Xinwei Long, Xuekai Zhu, Che Jiang, Yuchen Zhang, and 1 others. 2025. Ssr: Self-search reinforcement learning. *arXiv preprint arXiv:2508.10874*.

Tao Feng*, Xinke Jiang*, Xinyan Hu, Yonggang Zhang, Zhen Tao, Wentao Zhang, Boyang Liu, Wenhao Jiang, and Chao Wu. 2026. Trust within? seek beyond?

- knowledge boundary aware policy optimization for agentic search. In *ACL 2026*, pages 27664–27682.
- Jiaxuan Gao, Wei Fu, Minyang Xie, Shusheng Xu, Chuyi He, Zhiyu Mei, Banghua Zhu, and Yi Wu. 2025. [Beyond ten turns: Unlocking long-horizon agentic search with large-scale asynchronous rl](#). *Preprint*, arXiv:2508.07976.
- Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shitong Ma, Peiyi Wang, Xiao Bi, and 1 others. 2025. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*.
- Shibo Hao, Yi Gu, Haodi Ma, Joshua Jiahua Hong, Zhen Wang, Daisy Zhe Wang, and Zhiting Hu. 2023. Reasoning with language model is planning with world model. *arXiv preprint arXiv:2305.14992*.
- Hangfeng He, Hongming Zhang, and Dan Roth. 2022. [Rethinking with retrieval: Faithful large language model inference](#). *Preprint*, arXiv:2301.00303.
- Xanh Ho, Anh-Khoa Duong, Quoc-Huy Nguyen, and Suong Nguyen. 2020. Constructing a multi-hop qa dataset for comprehensive evaluation of reasoning steps. In *COLING*.
- Jakub Hoscilowicz, Bartosz Maj, Bartosz Kozakiewicz, Oleksii Tymoshchuk, and Artur Janicki. 2024. [Click-agent: Enhancing UI location capabilities of autonomous agents](#). *CoRR*, abs/2410.11872.
- Soyeong Jeong, Jinheon Baek, Sukmin Cho, Sung Ju Hwang, and Jong C Park. 2024. Adaptive-rag: Learning to adapt retrieval-augmented large language models through question complexity. *arXiv preprint arXiv:2403.14403*.
- Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. 2023a. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12):1–38.
- Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. 2023b. [Survey of hallucination in natural language generation](#). *ACM Comput. Surv.*, 55(12).
- Xinke Jiang, Yue Fang, Rihong Qiu, Haoyu Zhang, Yongxin Xu, Hao Chen, Wentao Zhang, Ruizhe Zhang, Yuchen Fang, Xu Chu, and 1 others. 2024a. Tc-rag: Turing-complete rag’s case study on medical llm systems. *arXiv preprint arXiv:2408.09199*.
- Xinke Jiang, Ruizhe Zhang, Yongxin Xu, Rihong Qiu, Yue Fang, Zhiyuan Wang, Jinyi Tang, Hongxin Ding, Xu Chu, Junfeng Zhao, and Yasha Wang. 2024b. [Hykge: A hypothesis knowledge graph enhanced framework for accurate and reliable medical llms responses](#). *Preprint*, arXiv:2312.15883.
- Bowen Jin, Hansi Zeng, Zhenrui Yue, Jinsung Yoon, Sercan Arik, Dong Wang, Hamed Zamani, and Jiawei Han. 2025. Search-R1: Training LLMs to reason and leverage search engines with reinforcement learning. *arXiv preprint arXiv:2503.09516*.
- Di Jin, Eileen Pan, Nassim Oufattole, Wei-Hung Weng, Hanyi Fang, and Peter Szolovits. 2020. What disease does this patient have? a large-scale open domain question answering dataset from medical exams. *arXiv preprint arXiv:2009.13081*.
- Mandar Joshi and 1 others. 2017. Triviaqa: A large scale distantly supervised challenge dataset for reading comprehension. In *ACL*.
- Nikhil Kandpal, Haikang Deng, Adam Roberts, Eric Wallace, and Colin Raffel. 2023. Large language models struggle to learn long-tail knowledge. In *International Conference on Machine Learning*, pages 15696–15707. PMLR.
- Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. 2020. [Scaling laws for neural language models](#). *Preprint*, arXiv:2001.08361.
- Urvashi Khandelwal, Omer Levy, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. 2020. [Generalization through memorization: Nearest neighbor language models](#). *Preprint*, arXiv:1911.00172.
- Satyapriya Krishna, Kalpesh Krishna, Anhad Mohananey, Steven Schwarcz, Adam Stambler, Shyam Upadhyay, and Manaal Faruqui. 2025. Fact, fetch, and reason: A unified evaluation of retrieval-augmented generation. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 4745–4759.
- Tom Kwiatkowski and 1 others. 2019. Natural questions: A benchmark for question answering research. *TACL*.
- Hanyu Lai, Xiao Liu, Iat Long Long, Shuntian Yao, Yuxuan Chen, Pengbo Shen, Hao Yu, Hanchen Zhang, Xiaohan Zhang, Yuxiao Dong, and 1 others. 2024. Autowebglm: A large language model-based web navigating agent. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 5295–5306.
- Myeonghwa Lee, Seonho An, and Min-Soo Kim. 2024. [PlanRAG: A plan-then-retrieval augmented generation for generative large language models as decision makers](#). In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 6537–6555, Mexico City, Mexico. Association for Computational Linguistics.

- Xiaoxi Li, Guanting Dong, Jiajie Jin, Yuyao Zhang, Yujia Zhou, Yutao Zhu, Peitian Zhang, and Zhicheng Dou. 2025. Search-o1: Agentic search-enhanced large reasoning models. *arXiv preprint arXiv:2501.05366*.
- Junlin Liu, Jiangwang Chen, Zixin Song, Shuaiyu Zhou, Chunji Lv, Hank Wu, Kailin Jiang, Jinyang Wu, Bohan Yu, and Chenxi Zhou. 2026. From proprietary to open-source: Bridging the distribution gap via multi-agent protocol distillation in agentic search. *arXiv preprint arXiv:2607.24280*.
- Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. 2023. [Lost in the middle: How language models use long contexts](#). *Preprint*, arXiv:2307.03172.
- Tianqi Liu, Yao Zhao, Rishabh Joshi, Misha Khalman, Mohammad Saleh, Peter J Liu, and Jialu Liu. Statistical rejection sampling improves preference optimization. In *The Twelfth International Conference on Learning Representations*.
- Yule Liu, Heyi Zhang, Jinyi Zheng, Zhen Sun, Zifan Peng, Jiaheng Wei, Tianshuo Cong, Yilong Yang, and Xinlei He. 2025. Auditing data membership in reinforcement learning with verifiable rewards. *arXiv preprint arXiv:2511.14045*.
- Chunji Lv, Yangguang Wei, Junlin Liu, Yang Gao, Ming Liu, Xinming Wang, Jinyang Wu, Guoren Wang, and Changsheng Li. 2026. Pcsd: Persistent consistency for self-distillation in agentic reinforcement learning. *arXiv preprint arXiv:2608.01837*.
- Liangbo Ning, Ziran Liang, Zhuohang Jiang, Haohao Qu, Yujuan Ding, Wenqi Fan, Xiao-yong Wei, Shanru Lin, Hui Liu, Philip S Yu, and 1 others. 2025. A survey of webagents: Towards next-generation ai agents for web automation with large foundation models. In *Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V. 2*, pages 6140–6150.
- OpenAI. 2025. [Openai-o1](#). Accessed: 2025-05-16.
- Akshara Prabhakar, Roshan Ram, Zixiang Chen, Silvio Savarese, Frank Wang, Caiming Xiong, Huan Wang, and Weiran Yao. 2025. Enterprise deep research: Steerable multi-agent deep research for enterprise analytics. *arXiv preprint arXiv:2510.17797*.
- Ofir Press, Muru Zhang, Sewon Min, Ludwig Schmidt, Noah A Smith, and Mike Lewis. 2022. Measuring and narrowing the compositionality gap in language models. *arXiv preprint arXiv:2210.03350*.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. 2023. Direct preference optimization: Your language model is secretly a reward model. *Advances in neural information processing systems*, 36:53728–53741.
- Daniel Russo and Benjamin Van Roy. 2014. [Learning to optimize via posterior sampling](#). *Preprint*, arXiv:1301.2609.
- Daniel Russo and Benjamin Van Roy. 2013. [Eluder dimension and the sample complexity of optimistic exploration](#). In *Advances in Neural Information Processing Systems*, volume 26. Curran Associates, Inc.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. 2017. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*.
- Zhihong Shao, Yeyun Gong, Yelong Shen, Minlie Huang, Nan Duan, and Weizhu Chen. 2023. Enhancing retrieval-augmented large language models with iterative retrieval-generation synergy. *arXiv preprint arXiv:2305.15294*.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, YK Li, Y Wu, and 1 others. 2024. Deepseek-math: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*.
- Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. Reflexion: Language agents with verbal reinforcement learning. *Advances in neural information processing systems*, 36:8634–8652.
- Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. 2021. [ALFWorld: Aligning Text and Embodied Environments for Interactive Learning](#). In *Proceedings of the International Conference on Learning Representations (ICLR)*.
- Harsh Trivedi, Niranjan Balasubramanian, Tushar Khot, and Ashish Sabharwal. 2022. Musique: Multi-hop questions via single-hop question composition. *Transactions of the Association for Computational Linguistics*, 10:539–554.
- Harsh Trivedi and 1 others. 2023. Interleaving retrieval with chain-of-thought reasoning for knowledge-intensive multi-step questions. In *ACL*.
- Minh Duc Vu, Han Wang, Jieshan Chen, Zhuang Li, Shengdong Zhao, Zhenchang Xing, and Chunyang Chen. 2024. Gptvoicetasker: Advancing multi-step mobile task efficiency through dynamic interface exploration and learning. In *Proceedings of the 37th Annual ACM Symposium on User Interface Software and Technology*, pages 1–17.
- Chenglong Wang, Yang Gan, Yifu Huo, Yongyu Mu, Qiaozhi He, Murun Yang, Bei Li, Tong Xiao, Chunliang Zhang, Tongran Liu, and Jingbo Zhu. 2025a. [Gram: A generative foundation reward model for reward generalization](#). *Preprint*, arXiv:2506.14175.
- Liang Wang, Haonan Chen, Nan Yang, Xiaolong Huang, Zhicheng Dou, and Furu Wei. 2025b. Chain-of-retrieval augmented generation. *arXiv preprint arXiv:2501.14342*.

- Shenzhi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen, Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen Yu, Gao Huang, and Junyang Lin. 2025c. [Beyond the 80/20 rule: High-entropy minority tokens drive effective reinforcement learning for llm reasoning](#). *Preprint*, arXiv:2506.01939.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, and 1 others. 2022a. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, and 1 others. 2022b. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837.
- Can Xie, Ruotong Pan, Xiangyu Wu, Zhang Yunfei, Jiayi Fu, Tingting Gao, and Guorui Zhou. 2026a. Unlocking exploration in rlvr: Uncertainty-aware advantage shaping for deeper reasoning. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 19057–19076.
- Can Xie, Yuyi Zhou, Wen Yang, Ziyi zhang, Siyao Song, Yingzhuo Deng, Shuo Ren, and Jiajun Zhang. 2026b. [Edge: Experience-distillation for guided exploration in agentic reinforcement learning](#). *Preprint*, arXiv:2608.21946.
- Wei Xiong, Jiarui Yao, Yuhui Xu, Bo Pang, Lei Wang, Doyen Sahoo, Junnan Li, Nan Jiang, Tong Zhang, Caiming Xiong, and 1 others. 2025. A minimalist approach to llm reasoning: from rejection sampling to reinforce. *arXiv preprint arXiv:2504.11343*.
- Weimin Xiong, Yifan Song, Xiutian Zhao, Wenhao Wu, Xun Wang, Ke Wang, Cheng Li, Wei Peng, and Sujian Li. 2024. Watch every step! llm agent learning via iterative step-level process refinement. *arXiv preprint arXiv:2406.11176*.
- Yongxin Xu, Xinke Jiang, Xu Chu, Rihong Qiu, Yujie Feng, Hongxin Ding, Junfeng Zhao, Yasha Wang, and Bing Xie. 2025a. Dearllm: Enhancing personalized healthcare via large language models-deduced feature correlations. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 941–949.
- Yongxin Xu, Ruizhe Zhang, Xinke Jiang, Yujie Feng, Yuzhen Xiao, Xinyu Ma, Runchuan Zhu, Xu Chu, Junfeng Zhao, and Yasha Wang. 2025b. Parenting: Optimizing knowledge selection of retrieval-augmented language models with parameter decoupling and tailored tuning. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11643–11662.
- Reda Yacoub and Dustin Axman. 2020. Probabilistic extension of precision, recall, and f1 score for more thorough evaluation of classification models. In *Proceedings of the first workshop on evaluation and comparison of NLP systems*, pages 79–91.
- An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Guanting Dong, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiaxin Yang, Jingren Zhou, Junyang Lin, and 25 others. 2024a. [Qwen2.5 technical report](#). *ArXiv*, abs/2412.15115.
- An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, and 1 others. 2024b. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*.
- Songhua Yang, Xinke Jiang, Hanjie Zhao, Wenxuan Zeng, Hongde Liu, and Yuxiang Jia. 2024c. Faima: Feature-aware in-context learning for multi-domain aspect-based sentiment analysis. In *COLING*.
- Zhibang Yang, Xinke Jiang, Rihong Qiu, Ruiqing Li, Yihang Zhang, Yue Fang, Yongxin Xu, Hongxin Ding, Xu Chu, Junfeng Zhao, and 1 others. 2026. Dfams: Dynamic-flow guided federated alignment based multi-prototype search. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 24728–24752.
- Zhilin Yang, Peng Qi, Saizheng Zhang, and 1 others. 2018. Hotpotqa: A dataset for diverse, explainable multi-hop question answering. In *EMNLP*.
- Shunyu Yao, Howard Chen, John Yang, and Karthik Narasimhan. 2022a. Webshop: Towards scalable real-world web interaction with grounded language agents. *Advances in Neural Information Processing Systems*, 35:20744–20757.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. 2022b. React: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*.
- Shenghao Ye, Yu Guo, Zhengheng Li, Shuangwu Chen, and Jian Yang. 2026. Rubric-guided process reward for stepwise model routing. *arXiv preprint arXiv:2605.29310*.
- Yixiao Zhou, Yang Li, Dongzhou Cheng, Hehe Fan, and Yu Cheng. 2026. Look inward to explore outward: Learning temperature policy from llm internal states via hierarchical rl. *arXiv preprint arXiv:2602.13035*.
- Zijian Zhou, Ao Qu, Zhaoxuan Wu, Sunghwan Kim, Alok Prakash, Daniela Rus, Jinhua Zhao, Bryan Kian Hsiang Low, and Paul Pu Liang. 2025. MEM1: Learning to synergize memory and reasoning for efficient long-horizon agents. *arXiv preprint arXiv:2506.15841*.

Appendix		M Broader Generalization	25
Contents		M.1 Medical Question Answering. . .	26
1 Introduction	1	M.2 Long-Form Research and Report Generation	26
2 AGENTICRAG-R1	3	M.3 Interactive Decision-Making Envi- ronments.	26
2.1 Action Modeling with Hierarchical Rewards	3	N Computational Resources and Software Environment	26
2.2 Information-Aware Trajectories Rejection	5	O The Use of Large Language Model	28
3 Experiments	6	A More Related Work	
3.1 Experimental Setup	6	A.1 Reinforcement Learning for LLMs	
3.2 Main Result Analysis	6	Reinforcement learning (RL) has emerged as a powerful paradigm for improving LLM alignment and reasoning, with recent work such as OpenAI- O1 (OpenAI, 2025) and DeepSeek-R1 (Guo et al., 2025) demonstrating strong performance on com- plex tasks (Ye et al., 2026; Zhou et al., 2026; Lv et al., 2026; Liu et al., 2025; Xie et al., 2026a,b). Methods like PPO (Schulman et al., 2017) have been widely adopted, but require multiple rounds of on-policy optimization, making them compu- tationally expensive and challenging to scale. To improve efficiency, alternatives such as Direct Pref- erence Optimization (DPO) (Rafailov et al., 2023) bypass reward modeling by directly optimizing preference loss, but suffer from off-policy bias. Group Relative Policy Optimization (GRPO) (Shao et al., 2024) offers a compromise by avoiding learned value networks and leveraging intra-group comparisons for relative rewards. Building on these, recent RL methods specifically designed for multi-step reasoning and tool interaction include AEPO (Dong et al., 2025a), which addresses en- tropy imbalance and rollout collapse in multi-round agents; ARPO (Dong et al., 2025b), which guides exploration of uncertain high-entropy tool calls us- ing advantage attribution; and MEM1 (Zhou et al., 2025), which synergizes memory and reasoning to support long-horizon multi-step interactions. De- spite these advances, existing RL approaches still face key challenges: coarse-grained action mod- eling, trajectory-level scalar rewards, limited exploration, cumulative errors in multi-step rea- soning, and the sparsity of training samples that require diverse reasoning behaviors.	
3.3 Component Analysis	7		
4 Related Works	9		
5 Conclusion and Future Work	9		
A More Related Work	14		
A.1 Reinforcement Learning for LLMs	14		
B Training Algorithm	15		
C Notations Table	15		
D Experiment Datasets	15		
D.1 Training Dataset.	15		
D.2 Testing set.	16		
E Baseline Implementation Details	17		
F Implementation Details	18		
G Prompt	18		
H Reward Analysis	19		
I Additional Experiments	19		
J Case Study	22		
K Agentic Hacking Theoretical Analysis	23		
K.1 Outcome-Only Reward Bias in Multi-Step Reasoning	23		
K.2 Bias Isolation via Pop-Based Stack and Context Masking	23		
L More Method Details: Pop-based Atten- tion Mask	25		

B Training Algorithm

Algorithm 1 details the full rollout engine with a maximum action budget B .

Algorithm 1 Dynamic Rollout with Multi-Action Reasoning and Stack-Based Memory

Require: query q , policy π_θ , retriever $\mathcal{R}$, budget T .

Ensure: Generated reasoning trajectory y

```

1: Initialize trajectory  $y \leftarrow \emptyset$  and memory stack  $\mathcal{M} \leftarrow \emptyset$ 
2: while  $t < T$  do
3:   Initialize action segment token  $x_t \leftarrow \emptyset$ 
4:   while  $x$  is not an action delimiter do
5:     Sample token  $x \sim \pi_\theta(\cdot \mid q, y, \mathcal{M})$ 
6:     Append  $x$  to  $x_t$ 
7:   end while
8:   Append  $x_t$  to trajectory  $y$ 
9:   Parse composed action  $a_t \leftarrow \text{parse}(x_t)$ 
10:  if  $a_t \in \{\text{Plan, Think, Search, Conclusion}\}$  then
11:    push( $\mathcal{M}$ , content of  $x_t$ )
12:  end if
13:  if  $a_t == \text{Search}$  then
14:    Extract query  $q_t^s \leftarrow \text{parse}(x_t)$ 
15:    Retrieve observation  $o_t \leftarrow \mathcal{R}(q_t^s)$ 
16:    Append  $\langle \text{Observation} \rangle o_t \langle / \text{Observation} \rangle$  to  $y$ 
17:    push( $\mathcal{M}$ ,  $\langle \text{Observation} \rangle o_t \langle / \text{Observation} \rangle$ )
18:    Modify Retrieval Loss Mask
19:  else if  $a_t == \text{Backtrack}$  then
20:    pop( $\mathcal{M}$ ), push( $\mathcal{M}$ , revised reasoning from  $x_t$ )
21:    Modify Pop-based Attention Mask
22:  else if  $a_t == \text{Summary}$  then
23:    pop( $\mathcal{M}$ ), push( $\mathcal{M}$ , summarized content)
24:    Modify Pop-based Attention Mask
25:  else if  $a_t == \text{Conclusion}$  then
26:    Return  $y$ 
27:  end if
28:   $t \leftarrow t + 1$ 
29: end while
30: Return  $y$ 

```

C Notations Table

Table 5 presents a comprehensive list of key notations and symbols used in the AGENTICRAG-R1 framework.

D Experiment Datasets

D.1 Training Dataset.

A central challenge in RL-based agent training lies in the limited difficulty of existing QA datasets. Effective supervision requires questions **that demand non-trivial reasoning and long-horizon, multi-step retrieval**. However, widely used multi-hop benchmarks such as HotpotQA and 2WikiMultiHopQA can often be solved with shallow reasoning and zero or only one retrieval step, making them insufficient for learning complex search and planning behaviors. As a result, models trained on such data **tend to overfit short-horizon, template-like reasoning patterns and generalize poorly to**

Symbol	Description
$\mathcal{S}$	System state capturing the current reasoning context
$\mathcal{A}$	Set of reasoning actions, consisting of primitive actions (push/pop) and composed actions (e.g., Plan, Search)
$\mathcal{M}$	Stack-based short-term memory of reasoning units
δ	State transition function over system state and memory
s_t	System state at reasoning step t
a_t	Reasoning action taken at step t
$\mathcal{M}_t$	Memory stack at step t
q	Input query sampled from the training dataset $\mathcal{D}$
y	Generated reasoning trajectory $(a_1, \dots, a_T)$
π_θ	Policy model parameterized by θ
π_{ref}	Frozen reference policy for KL regularization
K	Number of rollouts generated per input
$\mathcal{R}$	External retrieval module or search engine
$r_\phi(q, y)$	Scalar reward function evaluating trajectory y given input q
$x_i^{(k)}$	The i -th generated token in the k -th trajectory $y^{(k)}$.
R_G	The set of token-level rewards across all trajectories in the current rollout group
$r_i^{(k)}$	The token-level reward assigned to the token $x_i^{(k)}$.
$\hat{A}_{i,t}$	Group-relative advantage for trajectory i at step t
$\mathcal{J}(\theta)$	The Group Relative Policy Optimization (GRPO) objective.
$\tilde{z}_{i,t}$	Clipped importance sampling ratio for trajectory i at step t
$\mathbb{D}_{\text{KL}}(\cdot \parallel \cdot)$	Kullback–Leibler divergence between policies
β	KL regularization coefficient
$r_{\text{out}}^{(k)}$	Outcome-level reward evaluating final answer
$r_{\text{proc}}^{(k)}$	Process-level reward in the k -th trajectory
$\mathbb{I}_{\text{mask},i}^{(k)}$	Indicator function that equals 1 if the token $x_k^{(i)}$ is masked from policy optimization, and 0 otherwise.
$\mathbb{I}_{\text{act},i}^{(k)}$	Indicator function that equals 1 if the token $x_k^{(i)}$ corresponds to an action invocation, and 0 otherwise.
$a_{\text{conclusion}}^{(k)}$	The predicted final answer extracted from the $\langle \text{Conclusion} \rangle$ action in the k -th trajectory $y^{(k)}$
a_{gold}	The ground-truth answer used for evaluation
$r_{\text{fmt},i}^{(k)}$	Action format reward enforcing syntactic correctness at action level
$r_{\text{rag},i}^{(k)}$	Search-based reward evaluating retrieval relevance
$r_{\text{mem},i}^{(k)}$	Memory-aware reward for backtracking or summarization
$\mathbb{I}_{\text{syntax}}(a_i)$	Indicator for whether a_i has valid $\langle \text{ACTION} \rangle$ tags (1 = valid, 0 = invalid).
$\alpha_{\text{type}}(a_i)$	Action weight: 1.0 for CONCLUSION, 0.5 for other valid actions, 0 if invalid.
q_i^s	Search query issued at step i
o_i	Retrieved observation at step i
B	Mini-batch size (number of input prompts)
$\mathcal{B}$	Set of all rollouts in a batch
$\mathcal{B}^{\text{keep}}$	Rollout subset retained after global selection
$\mu^{(b)}$	Mean reward over rollouts for input $x^{(b)}$
$(\sigma^{(b)})^2$	Reward variance measuring input-level uncertainty
$\tilde{r}_i^{(b)}$	Information-aware score for rollout $\tau_i^{(b)}$
p	Fraction of top-ranked rollouts retained

Table 5: Key Notations Used in AGENTICRAG-R1

challenging agentic scenarios.

To address this limitation, we construct a curated training dataset following (Gao et al., 2025) by filtering instances from HotpotQA and 2Wiki that are both solvable and intrinsically require multi-step search. Specifically, for each candidate question, we prompt the training-free TC-RAG (Jiang et al., 2024a) to generate 10 independent trajectories under identical conditions, and discard questions that ❶ are answered incorrectly in all trials, ❷ achieve a correctness rate of 50% or higher, or ❸ can be answered correctly with at most a single retrieval step.

D.2 Testing set.

In our experiments, we evaluate our method on 9 widely used datasets spanning multi-hop reasoning, open-domain question answering, and real-world agent evaluation, including HotpotQA (Yang et al., 2018), 2WikiMultiHopQA (Ho et al., 2020), Bamboogle (Press et al., 2022), FRAMES (Krishna et al., 2025), MusiQue (Trivedi et al., 2022), NQ (Kwiatkowski et al., 2019) and TriviaQA (Joshi et al., 2017)). Dataset statistics, including the number of training, development, and test instances, are provided in Table 6:

- **2WikiMultiHopQA** (Ho et al., 2020) is a large-scale multi-hop question answering dataset constructed from Wikipedia and Wikidata. It contains 192,606 question-answer pairs grounded in a single database, with 154,878 training examples, 12,576 development examples, and 12,576 test examples. The dataset is designed to evaluate multi-hop reasoning that requires aggregating evidence across multiple documents.
- **HotpotQA** (Yang et al., 2018) is a large-scale dataset with 112,779 Wikipedia-based question-answer pairs requiring multi-hop reasoning over multiple documents. The dataset is divided into several subsets for training and evaluation. Specifically, the training set includes 18,089 single-hop questions (*train-easy*), 56,814 multi-hop questions (*train-medium*), and 15,661 hard multi-hop questions (*train-hard*). The development set contains 7,405 hard multi-hop questions, while the test set is further divided into two settings: *test-distractor* with 7,405 hard multi-hop questions and *test-fullwiki* with another 7,405 hard multi-hop questions. This division allows for a comprehensive evaluation of models under different conditions. HotpotQA also provides

sentence-level supporting facts required for reasoning, enabling models to improve performance and make explainable predictions.

- **Bamboogle** (Press et al., 2022) is a small, hand-crafted dataset with 125 questions designed to measure the extent to which a question answering system can answer varied compositional questions. It covers many different types of questions on various areas.
- **FRAMES** (Krishna et al., 2025) is a high-quality dataset designed to test LLMs’ factual responses, retrieval capabilities, and reasoning in generating final answers. It comprises 824 challenging multi-hop questions requiring integration of information from multiple sources.
- **MusiQue** (Trivedi et al., 2022) is a multihop question answering dataset constructed from Wikipedia, focusing on questions that require connected reasoning across multiple steps. It contains 19,938 training examples, 2,417 development examples, and 2,459 test examples. MusiQue is designed to enforce proper multihop reasoning by ensuring that each reasoning step critically relies on the output of the previous step.
- **Natural Questions (NQ)** (Kwiatkowski et al., 2019) is a large-scale question answering dataset based on real user queries issued. Each example pairs a question with a Wikipedia page and includes annotated long and short answers. The dataset contains 307,373 training examples, 7,830 development examples, and 7,842 test examples.
- **TriviaQA** (Joshi et al., 2017) is a large-scale reading comprehension dataset containing over 650K question-answer-evidence triples derived from 95K question-answer pairs authored by trivia enthusiasts. It includes evidence documents from Wikipedia and the Web, with an average of six documents per question. The dataset is divided into training (61,888 for Wiki, 76,496 for Web), development (7,993 for Wiki, 9,951 for Web), and test sets (7,701 for Wiki, 9,509 for Web). TriviaQA is designed to test models’ ability to handle complex questions with significant lexical and syntactic variability, requiring multi-sentence reasoning.

Dataset	Train	Dev	Test
2Wiki	154,878	12,576	12,576
HotpotQA	90,564	7,405	14,810
Bamboogle	0	0	125
FRAMES	0	0	824
MuSiQue	19,938	2,417	2,459
NQ	307,373	7,830	7,842
TriviaQA	138,384	17,944	17,210

Table 6: Overview of datasets used in experiments.

E Baseline Implementation Details

To comprehensively evaluate the effectiveness of the proposed method, we compare it against a diverse set of strong baselines covering different paradigms of reasoning, retrieval, and agentic interaction. The implementation details are described below.

❶ **No-RAG.** This category evaluates the model’s capability when no retrieval mechanism is involved, serving to assess how well the LLM performs using only its internal parametric knowledge.

- **No-RAG.** A non-retrieval baseline where the LLM directly generates responses based solely on its parametric knowledge, without access to any external information.
- **CoT (Chain-of-Thought) (Wei et al., 2022b).** Following prior work, this baseline prompts the model to generate explicit intermediate reasoning steps before producing the final answer.

❷ **Naive RAG.** Naive RAG methods incorporate retrieval in a heuristic or fixed manner without explicit reasoning–retrieval coordination, aiming to provide external knowledge through simple retrieval strategies.

- **FS-RAG (Fix Sentence RAG) (Trivedi et al., 2023).** FS-RAG performs retrieval at the sentence level, where each sentence in the input context is independently used as a retrieval query to obtain relevant external knowledge.
- **FL-RAG (Fix Length RAG) (Khandelwal et al., 2020).** FL-RAG invokes retrieval at fixed token intervals, triggering the retrieval module every (n) generated tokens using the preceding token window as the query.

❸ **Agentic RAG.** Agentic RAG methods rely on manually designed prompting strategies to interleave reasoning and retrieval, enabling the model to decide when and how to invoke external tools during inference. knowledge through simple retrieval strategies.

- **ReAct (Yao et al., 2022b).** ReAct interleaves reasoning and action steps, enabling the model to interact with external tools or environments during inference.
- **IRCOT (Trivedi et al., 2023).** IRCOT alternates between retrieval and Chain-of-Thought generation, where intermediate reasoning guides retrieval and retrieved evidence is used to refine reasoning.
- **TC-RAG (Jiang et al., 2024a).** TC-RAG introduces a Turing-complete RAG framework with a memory stack, allowing the model to dynamically control retrieval through push and pop.

❹ **RL-based Agentic RAG.** RL-based Agentic RAG approaches apply RL to retrieval-augmented settings, training models to learn when and how to invoke external retrieval tools during multi-step reasoning.

- **ReSearch (Chen et al., 2025).** ReSearch is trained using RL to jointly perform reasoning and search for multi-hop question answering, without requiring supervision over intermediate reasoning steps. We implement this baseline using Proximal Policy Optimization (PPO) (Schulman et al., 2017).
- **Search-R1 (Jin et al., 2025).** Search-R1 trains the LLM via RL to autonomously interact with a search engine during reasoning. It generates multi-round search queries and integrates retrieved information for problem solving. We train this baseline using Group Relative Policy Optimization (GRPO) (Shao et al., 2024).
- **AEPO (Dong et al., 2025a).** AEPO adopts an entropy-balanced policy optimization framework to stabilize training in high-entropy agentic settings by balancing entropy during rollout generation and policy updates.
- **ARPO (Dong et al., 2025b).** ARPO introduces an entropy-aware adaptive rollout strat-

egy that dynamically adjusts sampling at high-entropy decision points, encouraging diverse tool-use behaviors.

- **Mem1 (Zhou et al., 2025).** Mem1 is an end-to-end RL framework for long-horizon multi-turn tasks under constant memory constraints, maintaining a compact internal state for reasoning and memory consolidation.

F Implementation Details

We implement our method using a GRPO-based reinforcement learning framework. Training is performed with a batch size of 16 and a learning rate of 1×10^{-6} , without quantization or KV caching. We train for a single GRPO epoch consisting of 300 optimization steps. For each prompt, four stochastic generations are sampled with a temperature of 0.7, diverse sampling enabled, and a diversity penalty of 1.0. Each generation is limited to 2048 new tokens, with a maximum aggregated context length of 4096 tokens and up to ten iterative generation steps per episode. Rewards are computed at the token level. The GRPO optimizer follows a clipped objective with $\beta = 0.04$, $\mu = 2$, and $\epsilon = 0.1$. All hyperparameters are fixed across experiments, as training is conducted on a single curated dataset (Table 7).

Hyperparameter	Value
Batch Size	16
Learning Rate	1×10^{-6}
GRPO Epochs	1
Steps per Epoch	300
Number of Generations	10
Max New Tokens	2,048
Max Context Length	4,096
Max Generate Iterations	10
Temperature	0.7
β (GRPO)	0.04
μ (GRPO)	2
ϵ (GRPO)	0.1

Table 7: Implementation details and hyperparameter settings used in all experiments.

G Prompt

In this section, we provide the prompts used in our framework:

AGENTICRAG-R1 PROMPT TEMPLATE

Answer the following question to the best of your ability. You have access to the tools listed below.

[Insert tool_descriptions here]

The reasoning process and final answer should be wrapped in the `<ACTION>` and `</ACTION>` tags, respectively:

1. That is, "`<Think>` Write the reasoning process here `</Think>`";
2. "`<Conclusion>` Write the final answer here `</Conclusion>`";
3. Use the `<Backtrack>` tag to wrap your reflection results;
4. Use the `<Summary>` tag to wrap your summary results;
5. During the thinking process, ****if necessary, the assistant can perform a search**** to find uncertain knowledge, formatted as "`<Search>` search query. The search system will then provide the assistant with the retrieved information in the format "`<Observation>` ...search results... `</Observation>`".

Please think strictly according to the described process and use the following **Action Format**:

```
User_Query: ...
<Think> ... </Think>
<Search> ... </Search>
<Summary> ... </Summary>
<Backtrack> ... </Backtrack>
...(repeat <Think> / <Search> / <Summary> /
<Backtrack> as needed)
<Think> I now know the final answer. </Think>
<Conclusion> ... </Conclusion>
```

Begin!

GRM Relevance Evaluation Prompt

You are a relevance evaluator. Your task is to assess the semantic relevance of the retrieved documents with respect to the given question and task context, and output a relevance score between **0.0** and **1.0**, reported to **one decimal place** (e.g., 0.7).

Question: {Question}

Search Arguments: {Search Args}

Retrieved Documents: {Search Docs}

Scoring Guidelines:

- **1.0:** Highly relevant; directly answers the question and strongly supports reasoning.
- **0.7–0.9:** Mostly relevant and useful, with minor noise or redundancy.
- **0.4–0.6:** Partially relevant; indirect, incomplete, or weakly aligned.
- **0.1–0.3:** Largely irrelevant or misleading, with only marginal relevance.
- **0.0:** Completely irrelevant to the question.

Response Format:

Relevance Score:

GRM Memory Evaluation Prompt

You are a memory-operation evaluator. Your task is to assess whether the agent’s memory operation is reasonable and beneficial under the current question and stack context, and output a quality score between **0.0** and **1.0**, reported to **one decimal place** (e.g., 0.7).

Question: {Question}

Current Reasoning Context: {Reasoning Context}

Memory Stack Before Operation: {Stack Before}

Executed Memory Action: {Memory Action}

Memory Stack After Operation: {Stack After}

Scoring Guidelines:

- **1.0:** Highly reasonable and highly beneficial; the operation clearly improves the current reasoning state. For <Backtrack>, it removes misleading, redundant, or no-longer-useful memory while preserving useful information. For <Summary>, it effectively compresses prior context while preserving key facts needed for future reasoning.
- **0.7–0.9:** Mostly reasonable and useful, with only minor information loss, redundancy, or imperfect compression/removal.
- **0.4–0.6:** Partially reasonable; the intention is acceptable, but the operation is only weakly helpful, incomplete, or somewhat misaligned with the current reasoning needs.
- **0.1–0.3:** Largely unreasonable or weakly beneficial; the operation removes useful information, keeps misleading content, or produces an unhelpful summary.
- **0.0:** Completely unreasonable or harmful; the operation clearly damages the current reasoning state or is irrelevant to the task.

Evaluation Principles:

- Judge the operation with respect to the current question and reasoning state.
- Reward operations that help correct wrong branches, reduce distraction, or compress context without losing key evidence.
- Penalize operations that discard important facts, preserve misleading branches, or produce vague/inaccurate summaries.
- Focus on whether the operation is reasonable and beneficial at this step, not whether the final answer is ultimately correct.

Response Format:

Memory Operation Score:

H Reward Analysis

Our process rewards (r_{rag} and r_{mem}) are scored by a general-purpose LLM judge (GRM) via a fixed evaluation prompt, requiring no additional reward model training.

Semantic Reward Analysis. To assess whether semantic process evaluation is necessary, we replace the LLM judge with heuristic rewards:

ROUGE-L for retrieval relevance and a positive score for any valid memory operation. The average F1 score across seven benchmarks decreases from 33.55 to 29.40, and further drops to 19.94 when both r_{rag} and r_{mem} are removed, demonstrating the importance of semantic process supervision.

I Additional Experiments

Metric	TC-RAG	AGENTICRAG-R1
Avg.Time (s)	11.86	10.15
Avg.ResponseLen	838.14	880.89
Avg.Actions	10.14	11.38
Tool-Call Ratio (%)	91.70	97.50

Table 8: Statistical comparison of inference latency, response complexity, and tool-use behaviors between TC-RAG and AGENTICRAG-R1 on Qwen2.5-3B.

① Inference Behavior and Tool Usage. Table 8 compares AGENTICRAG-R1 with TC-RAG using Qwen2.5-3B. AGENTICRAG-R1 reduces average inference latency from 11.86 to 10.15 seconds (14.4%), while increasing the average number of actions by 12.2% and response length by 5.1%. Its tool-call rate increases from 91.7% to 97.5%, a difference of 5.8 percentage points. Thus, under this evaluation setting, the additional interactions do not result in higher measured average latency, although they produce more actions and slightly longer responses.

② Action-Specific Importance Analysis To analyze how different behaviors engage model parameters (Xu et al., 2025b; Chu et al., 2025b; Yang et al., 2026), we define a layer-wise importance score based on neuron activations. For each action token (including both the action token and its argument tokens), we extract its hidden representation at every transformer layer and count the number of neurons with positive activations after the ReLU nonlinearity:

$$\text{Importance}(l) = \sum \text{ReLU}(\text{embedding}_l > 0), \quad (9)$$

which measures intensity of neuron participation at layer l . Each embedding is 2048-dimensional and is visualized by averaging every 100 dimensions.

Importance scores are normalized within each action, as shown in the upper panels of Fig. 6. All statistics are computed on a challenging reasoning dataset collected via rejection sampling, retaining only correct-answer trajectories to avoid spurious




Figure 5: Training dynamics of AGENTICRAG-R1 during GRPO optimization on Qwen2.5-3B under different rejection ratios.

activations from erroneous reasoning. The top six heatmaps present the *normalized* importance profiles of six actions. At a coarse level, these actions exhibit broadly similar activation patterns, suggesting shared use of the model’s representation space. However, their importance distributions are not identical and encode action-specific differences in both magnitude and layer-wise allocation that are difficult to discern without direct comparison.

To make these distinctions explicit, the bottom row shows pairwise **differences computed from unnormalized importance scores**. Clear contrasts emerge between <Search> and <Think>, as well as between <Backtrack> and <Think>, indicating

substantial divergence in their underlying neuron activations. In contrast, the difference between <Summary> and <Backtrack> is relatively small, suggesting that these memory-related actions rely on highly overlapping neural resources. Overall, the results indicate that tool invocation (<Search>), internal reasoning (<Think>), and memory manipulation (<Backtrack>, <Summary>) correspond to distinct and largely decoupled activation patterns. **This finding supports explicit action-level decomposition and fine-grained optimization, rather than treating all behaviors as a single homogeneous reasoning process.**

Building on the action-level importance analy-




Figure 6: Neuron-wise importance analysis of different actions. Top two rows: normalized importance scores for individual actions. Bottom row: pairwise differences of unnormalized importance scores between actions.

sis above, we further examine activation patterns at a finer action–argument granularity. As shown in Fig. 7, we separately compute layer-wise importance scores for action tokens and their associated argument tokens using the same activation-sum metric as before, but aggregating activations across layers instead of tokens. While the overall trends remain consistent with the action-level results, clear divergences emerge across several contiguous layer ranges (highlighted in red), where different actions exhibit distinct layer-wise distributions for both their action tags and argument tokens. Importantly, although both components are consistently involved across actions, their relative contributions vary across layers in an action-dependent manner. This suggests that different actions are characterized by distinct layer allocation patterns between action tags and arguments, rather than differences in overall activation magnitude alone.

③ Action-Specific Optimization with Self-Search and Decoupling. Motivated by the observation that different actions rely on distinct and only partially overlapping neural resources, we investigate action-specific parameter optimization guided by model Shapley analysis. Specifically, we identify action-relevant parameters using SP-

R1 and decouple the optimization by assigning different subsets of parameters to the args layer and action layer, allowing each action to be updated independently. We evaluate this approach on MedQA under two search paradigms: External Search (ES), where both training and testing rely on external Wiki retrieval, and Self-Search (SS) (Fan et al., 2025), which generates candidate retrievals during training without external knowledge and uses external Wiki search only at test time. This self-supervised search encourages better generalization to unseen questions. Table 9 reports the results: incorporating SS alone improves accuracy over ES (from 65.5% to 67.5%), demonstrating its strong effect on generalization, while decoupled, action-specific optimization (Decoup) further boosts performance to 72.2% under SS and 68.5% under ES. These results highlight that SS significantly enhances generalization by exposing the model to diverse self-generated retrievals during training, and that Decoup effectively leverages action-specific parameter structure, enabling targeted optimization per action and yielding consistent gains across paradigms.

④ Training-Time Memory-Judge Token Cost. We quantify the token overhead of memory-




Figure 7: Layer-wise importance profiles of action tokens and their argument tokens.

operation evaluation during RL training. As shown in Table 10, each judge call consumes approximately 3,000–4,000 input tokens and 100–200 output tokens. Across a complete training run, memory-operation evaluation consumes approximately 3.5–4.0 million input tokens and 0.1–0.2 million output tokens. This overhead is incurred only during training and excludes retrieval-relevance judging.

⑤ Cross-Domain Rollout Reward Statistics.

To examine whether the trajectory rejection strategy remains effective across different domains and model scales, we report the mean and variance of rollout rewards. As shown in Table 11, meaningful rollout variance is preserved across all domains, confirming that the reward distributions do not collapse into indistinguishable ranges.

⑥ Scaling Analysis across Model Sizes. We compare average performance gains of AGENTICRAG-R1 over TC-RAG across three model scales. Smaller models benefit proportionally more from the framework, while larger models achieve the

Paradigm	Method	Acc (%)
Raw	Base	63.5
	React	65.5
ES	AGENTICRAG-R1+ ES	66.3
	AGENTICRAG-R1+ ES + Decoup	68.5
SS	AGENTICRAG-R1+ SS	67.5
	AGENTICRAG-R1+ SS+ Decoup	72.2

Table 9: Effect of action-specific (decoupled) parameter optimization under different search paradigms.

Component	Tokens per Call
Original context	2,500–3,000
Generated summary	300–500
Evaluation prompt and rubric	200–300
Total input	3,000–4,000
Output	100–200

Table 10: Estimated token consumption per memory-operation judge call.

best absolute performance:

- Qwen2.5-1.5B: 17.9 → 27.6 (+54%)
- Qwen2.5-3B: 22.2 → 33.5 (+50%)
- Qwen2.5-7B: 30.5 → 36.6 (+20%)

This pattern suggests that the structured action space and memory mechanisms are particularly effective at compensating for limited parametric capacity in smaller models.

J Case Study

We present a representative case study to illustrate the model’s ability to perform adaptive retrieval refinement and error correction. As shown in Table 12, the model initially issues broad retrieval queries that return irrelevant or noisy information, including unrelated medical content and off-topic encyclopedic entries. Instead of committing such noise to its internal state, the model explicitly invokes backtracking to revise retrieval keys and progressively narrow the search scope. Through multiple iterations of retrieval, inspection, and memory revision, the model correctly determines that *Battleship* is the 2012 military science fiction action film directed by Peter Berg in which Jason Snell worked. These cases highlight the effectiveness of fine-grained memory control and retrieval refinement in mitigating error accumulation during multi-step reasoning.

Domain	Qwen2.5-3B		Qwen2.5-7B	
	Mean	Var	Mean	Var
Math	0.44	0.07	0.46	0.09
Medical	0.37	0.10	0.53	0.12
General	0.41	0.12	0.65	0.11

Table 11: Rollout reward statistics across domains and model scales.

K Agentic Hacking Theoretical Analysis

K.1 Outcome-Only Reward Bias in Multi-Step Reasoning

Theorem 1 (Error Amplification under Outcome-Only Rewards). *Consider a multi-step reasoning trajectory $y = (a_1, a_2, \dots, a_T)$ generated by a policy $\pi_\theta(a_t | s_t)$ conditioned on an initial query q_0 . Let r_{out} denote an outcome-level reward assigned solely based on the correctness of the final answer. Under **outcome-only reward assignment without action-level reward decomposition**, each reasoning step a_t receives an identical implicit reward $\tilde{r}_t = \frac{r_{\text{out}}}{T}$, regardless of its semantic correctness¹. For any **erroneous or noisy step** a_e with true marginal contribution $\phi_e < 0$, the policy model parameter θ_k at step k update satisfies that:*

$$\theta_{k+1} = \theta_k + \mu \cdot \frac{r_{\text{out}}}{T} \cdot \nabla_\theta \log \pi_\theta(a_e | s_e), \quad (10)$$

where $\mu > 0$ is the learning rate, $\log \pi_\theta(a_e | s_e)$ is the output logits. Consequently, the generation probability of a_e increases whenever $r_{\text{out}} > 0$, irrespective of its negative contribution.

Proof. Under the RL estimator with outcome-only rewards, the policy gradient for a trajectory y is:

$$\nabla_\theta \mathbb{E}[R_y] = \mathbb{E} \left[r_{\text{out}} \sum_{t=1}^T \nabla_\theta \log \pi_\theta(a_t | s_t) \right]. \quad (11)$$

By linearity of expectation, this is equivalent to assigning each step an identical implicit reward $\tilde{r}_t = r_{\text{out}}/T$:

$$\nabla_\theta \mathbb{E}[R_y] = T \sum_{t=1}^T \mathbb{E}[\tilde{r}_t \cdot \nabla_\theta \log \pi_\theta(a_t | s_t)]. \quad (12)$$

This formulation implicitly assumes uniform contribution across steps. In practice, however, reasoning steps exhibit heterogeneous marginal contributions

¹Here we ignore the group-relative advantage normalization for simplicity.

$\{\phi_t\}_{t=1}^T$, where $\sum_{t=1}^T \phi_t = r_{\text{out}}$, and some ϕ_t may be negative due to hallucination, spurious inference, or retrieval noise. Thus, consider an erroneous step a_e with $\phi_e < 0$. The induced parameter update is

$$\Delta \theta = \mu \cdot \frac{r_{\text{out}}}{T} \cdot \nabla_\theta \log \pi_{\theta_k}(a_e | s_e). \quad (13)$$

Applying first-order Taylor expansion to step update:

$$\begin{aligned} \pi_{\theta_{k+1}}(a_e) &= \pi_{\theta_k}(a_e) + \nabla_\theta \pi_{\theta_k}(a_e) \cdot \Delta \theta + \mathcal{O}(\mu^2) \\ &= \pi_{\theta_k}(a_e) \left(1 + \mu \cdot \frac{r_{\text{out}}}{T} \left\| \nabla_\theta \log \pi_{\theta_k}(a_e | s_e) \right\|_2^2 \right) + \mathcal{O}(\mu^2). \end{aligned}$$

where we used the log-derivative identity $\nabla_\theta \pi = \pi \nabla_\theta \log \pi$. Since $r_{\text{out}} > 0$ for correct outcomes and gradient norms are non-negative, the bracketed term exceeds unity, implying

$$\pi_{\theta_{k+1}}(a_e) > \pi_{\theta_k}(a_e) \quad \text{regardless of } \phi_e < 0. \quad (14)$$

To characterize long-term accumulation, define the **step-level reward allocation error** as:

$$\varepsilon_e = \frac{r_{\text{out}}}{T} - \phi_e. \quad (15)$$

For erroneous steps, $\varepsilon_e > 0$ whenever $r_{\text{out}} > 0$. Over K training iterations, the cumulative bias satisfies

$$\mathbb{E}[\|B_e^{(K)}\|_2] \simeq \varepsilon_e \cdot K \cdot \mathbb{E}[\|\nabla_\theta \log \pi_\theta(a_e | s_e)\|_2], \quad (16)$$

demonstrating monotonic amplification of erroneous reasoning steps (**memory-aware actions**) in the absence of backtracking or summarizing. $\square$

K.2 Bias Isolation via Pop-Based Stack and Context Masking

Theorem 2 (Structural Isolation via Pop-Based Backtracking). *Consider a multi-step reasoning trajectory $y = (a_1, a_2, \dots, a_T)$ generated by a policy $\pi_\theta(a_t | s_t)$. Suppose that at step e , an erroneous action a_e is detected and removed via a pop-based stack mechanism, producing a pruned trajectory $\tilde{y}$. The removal is implemented by masking a_e from the conditioning context of all subsequent actions. Under this counterfactual trajectory $\tilde{y}$, the expected policy gradient update contains no contribution from a_e :*

$$\mathbb{E}[\nabla_\theta \log \pi_\theta(a_e | s_e) \cdot r_{\text{out}}] = 0. \quad (17)$$

Thus, erroneous reasoning steps do not receive outcome-level rewards and are not systematically reinforced.

Question	On which 2012 American military science fiction action film, directed by Peter Berg, did Jason Snell work?
Ground Truth	Battleship
Output	<Plan> First, I need to find the 2012 American military science fiction action film directed by Peter Berg. Then, I will search for the actors who worked on the film to identify Jason Snell. </Plan> <Search> [Wiki_RAG]: Peter Berg's films </Search> <Observation> '1. Gerald T. Olson is an American film producer best known for films such as House Party, Bloody Birthday and Dumb & Dumber. The San Francisco State alum relocated...' </Observation> <Backtrack> While the tool did provide some information about films directed by Peter Berg, it didn't specify the 2012 military science fiction action film I was looking for... Let me search for the 2012 military science fiction action film directed by Peter Berg. </Backtrack> <Search>[Wiki_RAG]: Search for detailed information about the 2012 military sci-fi action film Peter Berg</Search> <Observation> 'Wu Junyong (; born 1978) is a Chinese artist born in Putian, a city in Fujian province, China...' </Observation> <Backtrack> The tool didn't provide the information about the 2012 military science fiction action film directed by Peter Berg. I will now try searching for that specific film. </Backtrack> <Search> [Wiki_RAG]: Peter Berg movies 2012 </Search> <Observation> '1. House Party -1991 - drama 2. Bloody Birthday -1992 - horror 3. Dumb & Dumber -1994 - comedy 4. The Rundown -2003 - action 5. Patriots Day -2016 - drama The Hidden -1999 - horror Fast & Furious 6 -2013 - action' </Observation> <Backtrack> From the previous searches, no information about a 2012 military science fiction action film directed by Peter Berg and which Jason Snell worked on was retrieved. This information needs to be carefully re-evaluated and corrected. </Backtrack> <Search> [Wiki_RAG]: 2012 military science fiction action film directed by Peter Berg Jason Snell worked</Search> <Observation> '1. Peter Berg (born March 11, 1964) is an American director, producer, writer, and actor... the military science fiction war film Battleship (2012)...' </Observation> <Think> From the search result, I see that the 2012 movie directed by Peter Berg is "Battleship". However, the query mentions Jason Snell worked in it. Let me verify if Jason Snell indeed played a role in this film. </Think> <Search> [Wiki_RAG]: Jason Snell in Battleship</Search> <Observation> '1. Jason H. Snell is an American special... Indiana Jones and the Kingdom of the Battleship (2012)...' </Observation> <Summary> After verifying with the search results, it is clear that "Battleship" (2012) is the 2012 military science fiction action film directed by Peter Berg and in which Jason Snell worked. He played a role in this film.</Summary> <Conclusion> The 2012 military science fiction action film directed by Peter Berg in which Jason Snell worked is "Battleship".</Conclusion>

Table 12: A Case Study of a Wiki Problem on AGENTICRAG-R1.

Proof. Under outcome-only rewards, the standard advantage estimator is

$$\nabla_{\theta} \mathbb{E}[R_y] = \mathbb{E} \left[r_{\text{out}} \sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]. \quad (18)$$

Pop-based backtracking and context masking.

When a_e is removed, all subsequent actions are generated or evaluated under a masked context in which a_e is excluded. The resulting trajectory is

$$\tilde{y} = (a_1, \dots, a_{e-1}, a'_{e+1}, \dots, a'_T). \quad (19)$$

The attention mask ensures that a_e does not contribute to hidden states or logits involved in producing the final answer, effectively addressing the problem of reward hacking from outcome reward to be allocated the harmless popped actions.

Counterfactual independence. Under the masked context, the outcome distribution satisfies

$$p_{\theta}(a_T | a_1, \dots, a_{e-1}, \cancel{a_e}, \dots) = p_{\theta}(a_T | \tilde{y}), \quad (20)$$

implying that r_{out} depends only on $\tilde{y}$ and is conditionally independent of a_e .

Vanishing gradient contribution. The policy gradient under $\tilde{y}$ becomes

$$\nabla_{\theta} \mathbb{E}[R_{\tilde{y}}] = \mathbb{E} \left[r_{\text{out}} \sum_{t \in \tilde{y}} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]. \quad (21)$$

Since $a_e \notin \tilde{y}$, the corresponding score-function term never appears. Equivalently,

$$\mathbb{E}[\nabla_{\theta} \log \pi_{\theta}(a_e | s_e) \cdot r_{\text{out}}] = 0. \quad (22)$$

Therefore, erroneous reasoning steps are structurally isolated from outcome-level rewards. $\square$

Discussion. Rather than numerically redistributing rewards, the pop-based stack mechanism corrects outcome-only reward bias through *structural isolation*. By explicitly removing erroneous actions from the conditioning context via attention masking, policy gradients are evaluated on a counterfactual trajectory in which these actions never occurred, eliminating the error amplification effect identified in Section K.1.

L More Method Details: Pop-based Attention Mask

We extend the standard causal attention mechanism with a pop-based attention mask to handle dynamic context updates induced by <summary> and <backtrack> operations (<pop>). In autoregressive Transformers, causal attention ensures that each token at position (i) can only attend to tokens at positions ($j \leq i$). However, in our setting, causal ordering alone is insufficient, since tokens that appear earlier in the sequence may have been explicitly **removed from the active context** via a pop operation and should no longer be accessible.

Consider the example shown in Fig 2 Part 4, where the sequence follows an **Action-Summary** pattern: **Act. 1** $\rightarrow$ **Act. 2** $\rightarrow$ **Sum. 3** $\rightarrow$ **Act. 4**. The summary token **Sum. 3** aggregates the information of **Act. 2**, after which **Act. 2** is popped from the active memory. As a result, when generating **Act. 4**, the model should be allowed to attend to **Sum. 3**, but must be prevented from directly attending to **Act. 2**, even though **Act. 2** appears earlier in the sequence. Thus, the attention mask between **Act. 2** and **Act. 4** should be False. This constraint is indicated as **Mask after POP** in the attention matrix. Formally, let (M^{causal}) denote the standard causal mask:

$$M_{i,j}^{\text{causal}} = \begin{cases} 0, & j \leq i \\ -\infty, & j > i \end{cases} \quad (23)$$

We introduce an additional pop-based mask (M^{pop}), where a token (j) is masked if it has been popped before step (i):

$$M_{i,j}^{\text{pop}} = \begin{cases} -\infty, & \text{if } j \text{ is popped before } i \\ 0, & \text{otherwise} \end{cases} \quad (24)$$

The final attention mask is defined as:

$$M_{i,j} = M_{i,j}^{\text{causal}} + M_{i,j}^{\text{pop}}. \quad (25)$$

This formulation ensures that once an action has been summarized and popped, its information is only accessible through the corresponding summary token in all subsequent steps.

M Broader Generalization

To evaluate the effectiveness of AGENTICRAG-R1 across different application scenarios, we conduct additional experiments on MedQA, DeepResearch Bench, ALFWorld, and WebShop, covering medical question answering, long-form report generation, and interactive decision making.

Method	Format Acc (%)	Answer Acc (%)
Base	39%	84%
ReAct	56%	79%
AGENTICRAG-R1	92%	87%

Table 13: Results on the MedQA test set. The Qwen2.5-7B-Instruct model is evaluated using Qwen2.5-72B as the judge.

M.1 Medical Question Answering.

Recent work has demonstrated the potential of LLMs for domain-specific healthcare applications (Xu et al., 2025a; Jiang et al., 2024a; Feng* et al., 2026). To evaluate the generalization of AGENTICRAG-R1 in a domain-specific setting, we conduct experiments on MedQA (Jin et al., 2020). Table 13 reports the performance of different methods using Qwen2.5-72B as the judge. Compared with the base model, ReAct improves format accuracy from 39% to 56% but reduces answer accuracy from 84% to 79%. In contrast, AGENTICRAG-R1 achieves 92% format accuracy and 87% answer accuracy, outperforming both the base model and ReAct on both metrics. These results demonstrate that AGENTICRAG-R1 generalizes effectively to domain-specific medical question answering while improving format compliance without sacrificing answer quality.

M.2 Long-Form Research and Report Generation

Figure 8 presents the report-generation performance of AGENTICRAG-R1 on the 7B model, inferred with the query: “Give me a report of topic <RAG’s Future in Medical Domain>.” As shown, after multiple rounds of retrieval, the model produces a high-quality report accompanied by well-aligned citations. This result further demonstrates the strong generalization capability of our reinforcement learning-based agent training strategy.

To quantitatively evaluate the generalization of AGENTICRAG-R1 to long-horizon agentic scenarios beyond QA, we benchmark on DeepResearch Bench (Du et al., 2025), a multi-document report generation task that requires sustained multi-step retrieval, information aggregation, and structured composition. We evaluate using RACE (report quality) and FACT (factual consistency, including effective citation accuracy and claim-level accuracy). As shown in Table 14, AGENTICRAG-R1 substantially outperforms all baselines, including RAG, ReACT, and EDR (Prabhakar et al.,

2025), achieving the highest RACE score (43.60) and a significant margin on FACT metrics (+8.14 on Eff.c. and +8.56 on C.acc. over the best baseline).

Method	RACE	FACT (Eff.c.)	FACT (C.acc.)
RAG	36.68	2.93	10.86
ReACT	42.21	4.69	23.78
EDR	40.99	12.28	23.70
Ours	43.60	20.42	32.34

Table 14: Results on DeepResearch Bench for multi-document report generation.

M.3 Interactive Decision-Making Environments.

We further evaluate AGENTICRAG-R1 on interactive long-horizon decision making with ALFWorld (Shridhar et al., 2021) and WebShop (Yao et al., 2022a). These environments require agents to maintain state, incorporate feedback, and adapt strategies over extended trajectories. AGENTICRAG-R1 improves ALFWorld success rate from 36.5% to 51.1%, WebShop task score from 14.7 to 27.8, and WebShop success rate from 0.8% to 2.3%. Despite category-level variations on ALFWorld, these results demonstrate that AGENTICRAG-R1 generalizes beyond static QA to sequential interactive environments.

N Computational Resources and Software Environment

Experiments were carried out on a server equipped with two Intel Xeon Platinum 8488C processors, comprising 96 physical cores and 192 hardware threads in total. The system was configured with two NVIDIA A800 GPUs, each providing 80 GB of GPU memory, and 503 GB of system RAM. All experiments were conducted on Ubuntu 22.04.5 LTS. The software environment was based on Python 3.11.11, with package management handled through Conda (version 23.5.2). Model implementation and training were performed using PyTorch 2.6.0 with CUDA support, together with HuggingFace Transformers 4.51.3 and SpaCy 3.8.4. Unless otherwise specified, all models and libraries were used with their default parameter settings. The total training time for our method was approximately 21,582 seconds under the above hardware configuration. The computational cost of training, data preprocessing, and

Retrieval-Augmented Generation (RAG) in the Medical Field: Applications, Benefits, and Challenges

Key Points

- **Real-time Clinical Decision-Making:** RAG provides evidence-based recommendations, crucial for quick and accurate clinical decisions.
- **Medical Information Extraction:** Extracts and synthesizes medical data from patient records, literature, and guidelines, improving data accuracy.
- **Patient Diagnosis:** Assists in diagnosing patients by integrating relevant medical data with AI models.
- **Enhanced Accuracy and Transparency:** Combines model knowledge with external retrieval, reducing AI-generated misinformation.
- **Ethical and Privacy Concerns:** Significant challenges include privacy, bias, and data security issues that need to be addressed for responsible deployment.

Overview

Retrieval-Augmented Generation (RAG) is a cutting-edge technology that combines information retrieval and generative AI to enhance various aspects of healthcare. By leveraging both model knowledge and external data sources, RAG aims to provide more accurate, transparent, and up-to-date information, thereby improving clinical decision-making, patient diagnosis, and data management.

Overview

Retrieval-Augmented Generation (RAG) is a cutting-edge technology that combines information retrieval and generative AI to enhance various aspects of healthcare. By leveraging both model knowledge and external data sources, RAG aims to provide more accurate, transparent, and up-to-date information, thereby improving clinical decision-making, patient diagnosis, and data management.

Detailed Analysis

Applications of RAG in Medicine

1. **Real-time Clinical Decision-Making**
 - **Use Case:** RAG systems can provide healthcare professionals with evidence-based recommendations, enabling them to make quick and accurate decisions.
 - **Source:** [A Scoping Review of Retrieval-Augmented Generation—Ethics, Bias](#)
2. **Medical Information Extraction**
 - **Use Case:** RAG can extract and synthesize medical data from various sources, such as patient records, research literature, and clinical guidelines, improving the accuracy and completeness of medical information.
 - **Source:** [Enhancing medical AI with retrieval-augmented generation: A mini](#)
3. **Patient Diagnosis**
 - **Use Case:** RAG can assist in diagnosing patients by integrating relevant medical data with AI models, providing a more comprehensive and accurate diagnosis.
 - **Source:** [Enhancing medical AI with retrieval-augmented generation: A mini](#)
4. **Data Accessibility**
 - **Use Case:** RAG on FHIR (Fast Healthcare Interoperability Resources) makes healthcare data more accessible by retrieving real-time data based on specific questions.

(a) Report Part (1/2)

- **Source:** [A Scoping Review of Retrieval-Augmented Generation—Ethics, Bias](#)
3. **Reduction of Hallucinations**
 - **Benefit:** Grounding AI responses in external knowledge significantly reduces the occurrence of incorrect or irrelevant information.
 - **Source:** [Retrieval augmented generation for large language models and...](#)
 4. **Up-to-date Information**
 - **Benefit:** RAG can incorporate the latest clinical data and guidelines, ensuring that the information provided is current and relevant.
 - **Source:** [Retrieval augmented generation for large language models and...](#)

Challenges of RAG in Medicine

1. **Lack of Standardized Evaluation Frameworks**
 - **Challenge:** There is a need for standardized frameworks to assess the performance and reliability of RAG systems.
 - **Source:** [A Scoping Review of Retrieval-Augmented Generation—Ethics, Bias](#)
2. **Ethical Concerns**
 - **Challenge:** Privacy and bias must be addressed to ensure the responsible deployment of RAG in healthcare.
 - **Source:** [A Scoping Review of Retrieval-Augmented Generation—Ethics, Bias](#)
3. **Complexity of Retrieval Mechanisms**
 - **Challenge:** Optimizing the retrieval mechanism to minimize irrelevant information is complex and requires advanced techniques.
 - **Source:** [A Scoping Review of Retrieval-Augmented Generation—Ethics, Bias](#)
4. **Missing Content and Hallucinations**

4. **Missing Content and Hallucinations**
 - **Challenge:** RAG may still provide incorrect answers if relevant information is not available in the knowledge base.
 - **Source:** [Top 7 Challenges with Retrieval-Augmented Generation - Valprovia](#)
5. **Performance and Scalability**
 - **Challenge:** Handling large datasets and high query volumes efficiently is a significant challenge.
 - **Source:** [Top Problems with RAG systems and ways to mitigate them](#)
6. **Security Concerns**
 - **Challenge:** Data leakage and privacy risks are significant concerns that need to be addressed.
 - **Source:** [Retrieval-augmented generation for generative artificial intelligence...](#)

Key Citations

- [A Scoping Review of Retrieval-Augmented Generation—Ethics, Bias](#)
- [Enhancing medical AI with retrieval-augmented generation: A mini](#)
- [Generative AI in Healthcare: Use Cases, Benefits, and Challenges](#)
- [Retrieval augmented generation for large language models and...](#)
- [Top 7 Challenges with Retrieval-Augmented Generation - Valprovia](#)
- [Top Problems with RAG systems and ways to mitigate them](#)
- [Retrieval-augmented generation for generative artificial intelligence...](#)

(b) Report Part (2/2)

Figure 8: Report Generation for query "Give me a report of topic <RAG's Future in Medical Domain>"

Method	ALFWorld							WebShop	
	Pick	Look	Clean	Heat	Cool	Pick2	Overall	Score	SR
Prompt	45.8	48.4	34.5	41.0	19.6	31.7	36.5	14.7	0.8
AGENTICRAG-R1	84.7	80.6	24.1	33.3	37.0	51.2	51.1	27.8	2.3

Table 15: Generalization of AGENTICRAG-R1 to interactive decision-making environments. ALFWorld reports task-specific and overall success rates (%); WebShop reports the official task score and success rate (SR, %).

inference varied depending on the experimental configuration and dataset scale, and all reported results were obtained under the same hardware and software setup.

O The Use of Large Language Model

In this study, generative AI software tools were used solely as auxiliary support for language refinement, programming-related assistance, and figure visualization. Specifically, large language models were employed to improve grammatical accuracy, readability, and overall writing quality, as well as to provide general coding suggestions or debugging guidance. In addition, AI-based image generation tools were used to enhance the visual appearance of figure 1; however, the figure design, initial draft, and final verification were entirely performed by the authors. All AI-assisted outputs were carefully reviewed and validated by the authors prior to inclusion. The research conception, methodological design, and analysis of experimental results were conducted exclusively by the authors, and generative AI tools did not contribute to the formulation of research ideas or the derivation of conclusions.